

March 24th, 2021

AI+BD ML Lab. Day 3

MLP & Regularization (Batch Norm. / Dropout)

YoungIn Kim

<youngkim21@postech.edu>



Contents

1. Today's Goals

2. MLP (Multi Layer Perceptron)

- ✦ Make model with nn.Module Class
- ✦ Make more deep!

3. Regularization

- ✦ Look inside of network
- ✦ Batch Normalization
- ✦ Dropout



Before we start,
let's remind yesterday's code

Base .ipynb :

<https://git.io/aibd-mlp-3>

Make model with nn.Module class

4



1. Library Importation & Device Preparation

```
1  # Library Importation
2  import matplotlib.pyplot as plt
3  import numpy as np
4  import seaborn as sns
5  import time
6  import torch
7  import torch.nn as nn
8  import torch.nn.functional as F
9
10 from IPython.display import clear_output
11 from multiprocessing import cpu_count
12 from sklearn.metrics import confusion_matrix
13 from torch.optim import SGD
14 from torch.utils.data import DataLoader, random_split
15 from torchvision.datasets import MNIST
16 from torchvision.transforms import ToTensor
17
18 MNIST.resources = [
19     (
20         'https://ossci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte.gz',
21         'f68b3c2dcbeaaa9fbdd348bbdeb94873'
22     ), (
23         'https://ossci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte.gz',
24         'd53e105ee54ea40749a09fcbcd1e9432'
25     ), (
26         'https://ossci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz',
27         '9fb629c4189551a2d022fa330f9573f3'
28     ), (
29         'https://ossci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz',
30         'ec29112dd5afa0611ce80d1b7f02629c'
31     )
32 ]
33
34 # Device Preparation
35 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
36 print(f'{"CPU" if device == "cpu" else "GPU"} will be used in training/validation.')
```



2. Hyper-Parameters

```
1  # Data Loader
2  batch_size = 32
3
4  # Model
5  hidden_layer = 200
6
7  # Learning
8  logging_disfig = True
9  maximum_epoch = 25
10 learning_rate = 0.1
```



3. Data Load & Preprocessing

```
1 # Load dataset into python variable
2 train_data = MNIST("./", train=True, transform=ToTensor(), target_transform=None, download=True)
3 train_data, valid_data = random_split(train_data, [54000, 6000])
4 test_data = MNIST("./", train=False, transform=ToTensor(), target_transform=None, download=True)
5
6 # Check the data
7 print('===== Check the data =====\n')
8 print(f'Train dataset length = {len(train_data)}')
9 print(f'Valid dataset length = {len(valid_data)}')
10 print(f'Test dataset length = {len(test_data)}\n')
11
12 train_0_x, train_0_y = train_data[0]
13 print(f'Content of Y (Label, type={type(train_0_y)}) = {train_0_y}')
14 print(f'Shape of X (Data, type={type(train_0_x)}) = {train_0_x.shape}')
15 plt.figure(1)
16 plt.imshow(train_0_x.squeeze())
17 plt.title(f'train_0_x ({train_0_x.squeeze().shape})')
18 plt.show()
19
20 # Create data loader
21 train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True, pin_memory=True,
22                           drop_last=True)
23 valid_loader = DataLoader(valid_data, batch_size=len(valid_data), pin_memory=True)
24 test_loader = DataLoader(test_data, batch_size=len(test_data), pin_memory=True)
25
26 # Examine the data loader
27 print('===== Check the data loader =====\n')
28 train_enumerator = enumerate(train_loader)
29 ex_batch_idx, (ex_data, ex_label) = next(train_enumerator)
30 print(f'Idx: {ex_batch_idx} / X.shape = {ex_data.shape} / Y.shape = {ex_label.shape}\n')
31 print(f'Y[0:{batch_size}] = {ex_label}')
32
33 preview_index = 0
34 plt.figure(2)
35 plt.imshow(ex_data[preview_index, 0, :, :])
36 plt.title(f'Batch example data [{preview_index}, label={ex_label[preview_index]}]')
37 plt.show()
```



4. Function Definitions

```
1 # Model
2 def init_model(_net):
3     global net, loss_fn, optim
4     net = _net.to(device)
5     loss_fn = nn.CrossEntropyLoss()
6     optim = SGD(net.parameters(), lr=learning_rate)
7
8 # Epoch
9 def init_epoch():
10     global epoch_cnt
11     epoch_cnt = 0
12
```



4. Function Definitions (continue)



IMPORTANT!

```
14 def epoch(data_loader):
15     # One epoch : gets data_loader as input and returns loss / accuracy, and
16     #                 last prediction value / its label(truth) value for future use
17     global epoch_cnt
18     iter_loss, iter_acc = [], []
19
20     last_out, last_label = None, None
21     last_grad_performed = False
22
23     # Mini-batch iterations
24     for _data, _label in data_loader:
25         data, label = _data.view([len(_data), -1]).to(device), _label.to(device)
26
27         # 1. Feed-forward
28         onehot_out = net(data)
29
30         # 2. Calculate accuracy
31         _, out = torch.max(onehot_out, 1)
32         acc_partial = (out == label).float().sum()
33         acc_partial = acc_partial / len(label)
34         iter_acc.append(acc_partial.item())
35
36         # 3. Calculate loss
37         loss = loss_fn(onehot_out, label)
38         iter_loss.append(loss.item())
39
40         # 4. Backward propagation if not in 'torch.no_grad()'
41         if onehot_out.requires_grad:
42             optim.zero_grad()
43             loss.backward()
44             optim.step()
45             last_grad_performed = True
46
47         # 5. Save current iteration data for future use
48         last_out = out.cpu().detach()
49         last_label = _label
```




4. Function Definitions (continue)

```
51     # Up epoch count if backward propagation is done
52     if last_grad_performed:
53         epoch_cnt += 1
54
55     return np.average(iter_loss), np.average(iter_acc), last_out, last_label
56
57
58 def epoch_not_finished():
59     # For now, let's repeat training fixed times, e.g. 25 times.
60     # We will learn how to determine training stop or continue later.
61     return epoch_cnt < maximum_epoch
```



5. Model Architectures (before)

```
1 # before
2
3 net = nn.Linear(1, 1)
4
5 net = nn.Sequential(
6     nn.Linear(len(train_0_x.view([-1])), hidden_layer, bias=False),
7     nn.ReLU(),
8     nn.Linear(hidden_layer, 10, bias=False)
9 ).to(device)
```



6. Training Iteration & Test Result

```
1 # Training Initialization
2 init_model(OneLayerModel())
3 init_epoch()
4 init_log()
5
6 # Training Iteration
7 while epoch_not_finished():
8     start_time = time.time()
9     tloss, tacc, _, _ = epoch(train_loader)
10    end_time = time.time()
11    time_taken = end_time - start_time
12    record_train_log(tloss, tacc, time_taken)
13    with torch.no_grad():
14        vloss, vacc, _, _ = epoch(valid_loader)
15        record_valid_log(vloss, vacc)
16    print_log()
17
18 print('\n Training completed!')
19
20 # Accuracy for test dataset
21 with torch.no_grad():
22     test_loss, test_acc, test_out, test_label = epoch(test_loader)
23     print('\n===== Test Result =====\n')
24     print(f'Test accuracy = {test_acc}\nTest loss = {test_loss}')
```

```
1 # Model
2 def init_model(_net):
3     global net, loss_fn, optim
4     net = _net.to(device)
5     loss_fn = nn.CrossEntropyLoss()
6     optim = SGD(net.parameters(), lr=learning_rate)
7
```



5. Model Architectures (Back again)

```
1 # before
2
3 net = nn.Linear(1, 1)
4
```



```
1 class LinearModel(nn.Module): # nn.Module 상속
2     def __init__(self): # 속성값 초기화
3         super(LinearModel, self).__init__()
4         self.linear = nn.Linear(1, 1)
5
6     def forward(self, x): # 학습데이터를 입력받아서 forward 연산 진행
7         return self.linear(x)
8
9 model = LinearModel()
```



5. Model Architectures (Back again)

```
5 net = nn.Sequential(  
6     nn.Linear(len(train_0_x.view([-1])), hidden_layer, bias=False),  
7     nn.ReLU(),  
8     nn.Linear(hidden_layer, 10, bias=False)  
9 ).to(device)
```




```
1 class OneLayerModel(nn.Module):  
2     def __init__(self):  
3         super(OneLayerModel, self).__init__()  
4  
5         self.fc1 = nn.Linear(len(train_0_x.view([-1])), hidden_layer, bias=False)  
6         self.act = nn.ReLU()  
7         self.fc2 = nn.Linear(hidden_layer, 10, bias=False)  
8  
9     def forward(self, x):  
10        out = self.fc1(x)  
11        hidden = self.act(out)  
12        onehot_out = self.fc2(hidden)  
13  
14        return onehot_out
```



6. Training Iteration & Test Result

```
1 class OneLayerModel(nn.Module):
2     def __init__(self):
3         super(OneLayerModel, self).__init__()
4
5         self.fc1 = nn.Linear(len(train_0_x.view([-1])), hidden_layer, bias=False)
6         self.act = nn.ReLU()
7         self.fc2 = nn.Linear(hidden_layer, 10, bias=False)
8
9     def forward(self, x):
10        out = self.fc1(x)
11        hidden = self.act(out)
12        onehot_out = self.fc2(hidden)
13
14        return onehot_out
```



Iter: 25 >> T_loss 0.00487 T_acc 0.99972 V_loss 0.05918 V_acc 0.98233 3.036s

===== Test Result =====

Test accuracy = 0.9801999926567078
Test loss = 0.06666535884141922



6-1. Training Iteration & Test Result (MLP)

```
1 class MLP(nn.Module):
2     def __init__(self, in_features, out_features):
3         super(MLP, self).__init__()
4
5         self.hidden_layer1 = 165
6         self.hidden_layer2 = 165
7
8         self.fc1 = nn.Linear(in_features, self.hidden_layer1)
9         self.act1 = nn.ReLU()
10        self.fc2 = nn.Linear(self.hidden_layer1, self.hidden_layer2)
11        self.act2 = nn.ReLU()
12        self.fc3 = nn.Linear(self.hidden_layer2, out_features)
13
14    def forward(self, x):
15        hidden1 = self.act1(self.fc1(x))  ## self.act nn.ReLU() -> instance of class
16        hidden2 = F.relu(self.fc2(hidden1))  ## F.relu -> function
17
18        onehot_out = self.fc3(hidden2)
19
20    return onehot_out
```

```
7 import torch.nn as nn
8 import torch.nn.functional as F
```

```
1 # Training Initialization
2 init_model(MLP(len(train_0_x.view([-1])), 10))
3 init_epoch()
4 init_log()
5
```


Make more Deep!

16



6-1. Training Iteration & Test Result (MLP)

```
1 class MLP(nn.Module):
2     def __init__(self, in_features, out_features):
3         super(MLP, self).__init__()
4
5         self.hidden_layer1 = 165
6         self.hidden_layer2 = 165
7
8         self.fc1 = nn.Linear(in_features, self.hidden_layer1)
9         self.act1 = nn.ReLU()
10        self.fc2 = nn.Linear(self.hidden_layer1, self.hidden_layer2)
11        self.act2 = nn.ReLU()
12        self.fc3 = nn.Linear(self.hidden_layer2, out_features)
13
14    def forward(self, x):
15        hidden1 = self.act(self.fc1(x))  ## self.act nn.ReLU() -> instance of class
16        hidden2 = F.relu(self.fc2(hidden1))  ## F.relu -> function
17
18        onehot_out = self.fc3(hidden2)
19
20    return onehot_out
```



Iter: 25 >> T_loss 0.00047 T_acc 1.00000 V_loss 0.07750 V_acc 0.98267 ⌚ 3.298s

Test accuracy = 0.982699990272522

Test loss = 0.07591626048088074

Batch Normalization – Look inside of network

17



```
1 def plot_inner_dist():
2     for epoch in range(10):
3         net.train()
4         for _data, _label in train_loader:
5             data, label = _data.view([len(_data), -1]).to(device), _label.to(device)
6
7             # Feed-forward
8             _, _, _, onehot_out = net(data)
9             loss = loss_fn(onehot_out, label)
10
11            # Backward propagation
12            optim.zero_grad()
13            loss.backward()
14            optim.step()
15
16        net.eval()
17        with torch.no_grad():
18            for _data, _label in test_loader:
19                data, label = _data.view([len(_data), -1]).to(device), _label.to(device)
20
21                # Feed-forward
22                o1, h1, o2, h2, onehot_out = net(data)
23
24                _, out = torch.max(onehot_out, 1)
25                acc_test = (out == label).float().sum()
26                acc_test = acc_test / len(label)
27
28            # plot inner distribution
29            o1, h1, o2, h2 = o1.cpu().detach().numpy(), h1.cpu().detach().numpy(), o2.cpu().detach().numpy(),
30            fig, axs = plt.subplots(2, 3, figsize=(10, 7), sharex='col')
31            axs[0, 0].hist(o1.reshape(-1))
32            axs[0, 1].hist(h1.reshape(-1))
33            axs[0, 2].scatter(o1[0], h1[0])
34            axs[1, 0].hist(o2.reshape(-1))
35            axs[1, 1].hist(h2.reshape(-1))
36            axs[1, 2].scatter(o2[0], h2[0])
37
38        clear_output(wait=True)
39        print(f'\n### Epoch: {epoch+1} / Accuracy: {acc_test} ###')
40        plt.show()
```

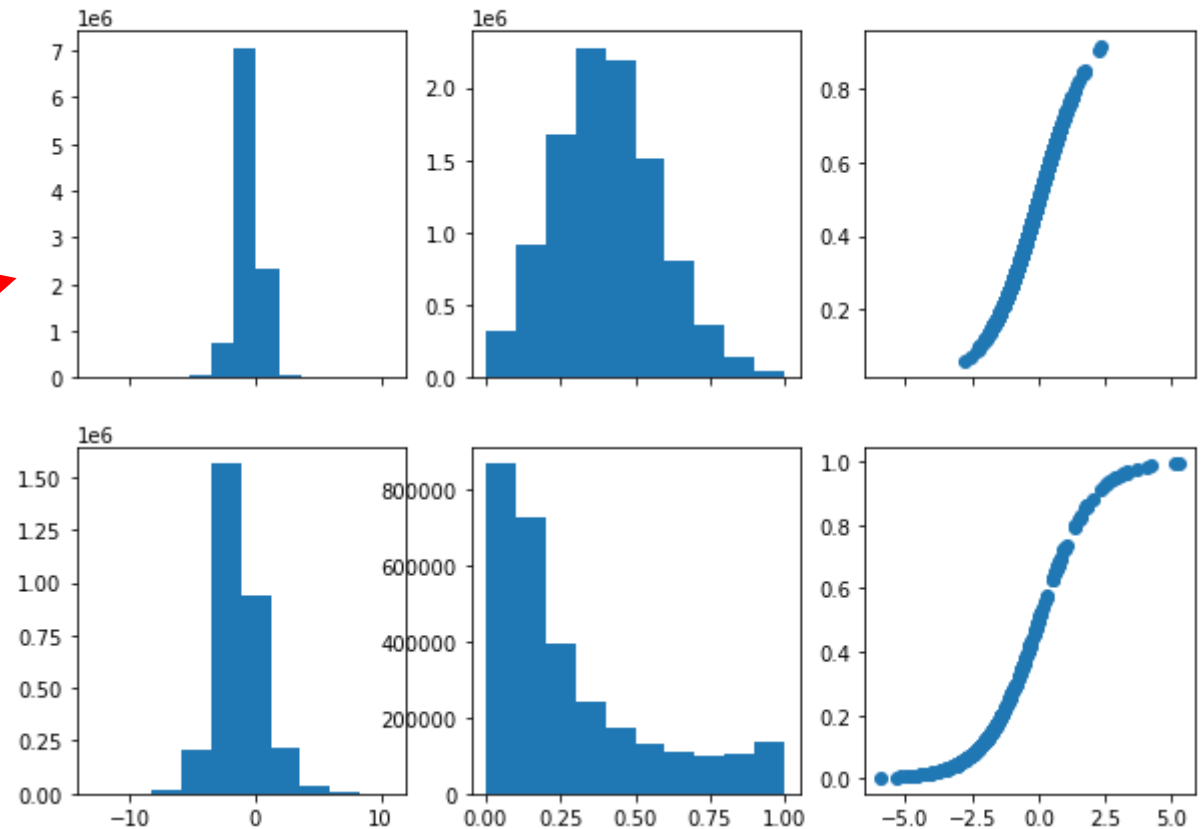
Batch Normalization – Look inside of network

18



```
1 class MLP_Sigmoid(nn.Module):
2     def __init__(self, in_features, out_features):
3         super(MLP_Sigmoid, self).__init__()
4
5         self.hidden_layer1 = 1024
6         self.hidden_layer2 = 300
7
8         self.fc1 = nn.Linear(in_features, self.hidden_layer1)
9         self.act1 = nn.Sigmoid()
10        self.fc2 = nn.Linear(self.hidden_layer1, self.hidden_layer2)
11        self.act2 = nn.Sigmoid()
12        self.fc3 = nn.Linear(self.hidden_layer2, out_features)
13
14
15    def forward(self, x):
16        output1 = self.fc1(x)
17        hidden1 = self.act1(output1)
18
19        output2 = self.fc2(hidden1)
20        hidden2 = self.act2(output2)
21
22        onehot_out = self.fc3(hidden2)
23
24        return output1, hidden1, output2, hidden2, onehot_out
25
26
27    init_model(MLP_Sigmoid(len(train_0_x.view([-1])), 10).to(device))
28    plot_inner_dist()
```

Epoch: 10 / Accuracy: 0.9307000041007996




Batch Normalization – Look inside of network

19



```
1 class MLPTanh(nn.Module):
2     def __init__(self, in_features, out_features):
3         super(MLPTanh, self).__init__()
4
5         self.hidden_layer1 = 1024
6         self.hidden_layer2 = 300
7
8         self.fc1 = nn.Linear(in_features, self.hidden_layer1)
9         self.act1 = nn.Tanh()
10        self.fc2 = nn.Linear(self.hidden_layer1, self.hidden_layer2)
11        self.act2 = nn.Tanh()
12        self.fc3 = nn.Linear(self.hidden_layer2, out_features)
13
14    def forward(self, x):
15        output1 = self.fc1(x)
16        hidden1 = self.act1(output1)
17
18        output2 = self.fc2(hidden1)
19        hidden2 = self.act2(output2)
20
21        onehot_out = self.fc3(hidden2)
22
23        return output1, hidden1, output2, hidden2, onehot_out
24
25 init_model(MLPTanh(len(train_0_x.view([-1])), 10).to(device))
26 plot_inner_dist()
```

Epoch: 10 / Accuracy: 0.9781000018119812







Batch Normalization – Look inside of network

20



Internal Covariate Shift



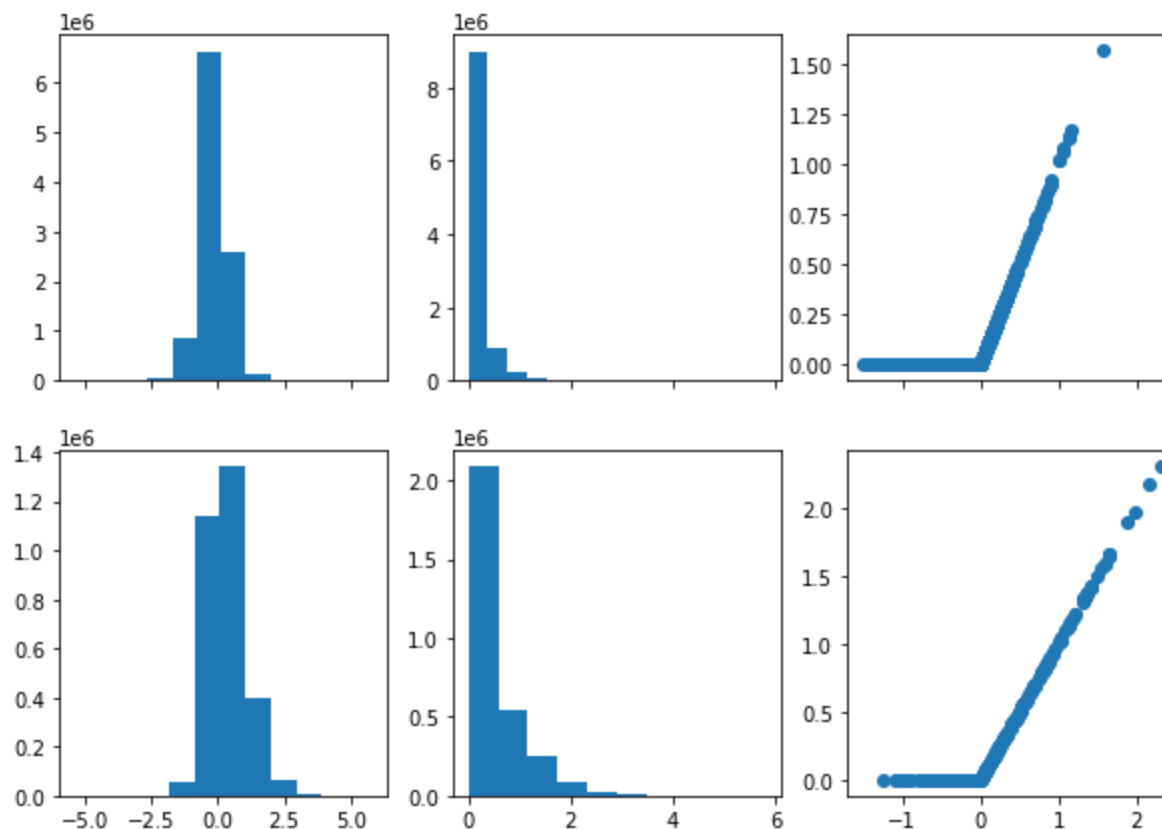
Batch Normalization – Look inside of network

21



```
1 class MLPReLU(nn.Module):
2     def __init__(self, in_features, out_features):
3         super(MLPReLU, self).__init__()
4
5         self.hidden_layer1 = 1024
6         self.hidden_layer2 = 300
7
8         self.fc1 = nn.Linear(in_features, self.hidden_layer1)
9         self.act1 = nn.ReLU()
10        self.fc2 = nn.Linear(self.hidden_layer1, self.hidden_layer2)
11        self.act2 = nn.ReLU()
12        self.fc3 = nn.Linear(self.hidden_layer2, out_features)
13
14    def forward(self, x):
15        output1 = self.fc1(x)
16        hidden1 = self.act1(output1)
17
18        output2 = self.fc2(hidden1)
19        hidden2 = self.act2(output2)
20
21        onehot_out = self.fc3(hidden2)
22
23        return output1, hidden1, output2, hidden2, onehot_out
24
25 init_model(MLPReLU(len(train_0_x.view([-1])), 10).to(device))
26 plot_inner_dist()
```

Epoch: 10 / Accuracy: 0.9822999835014343



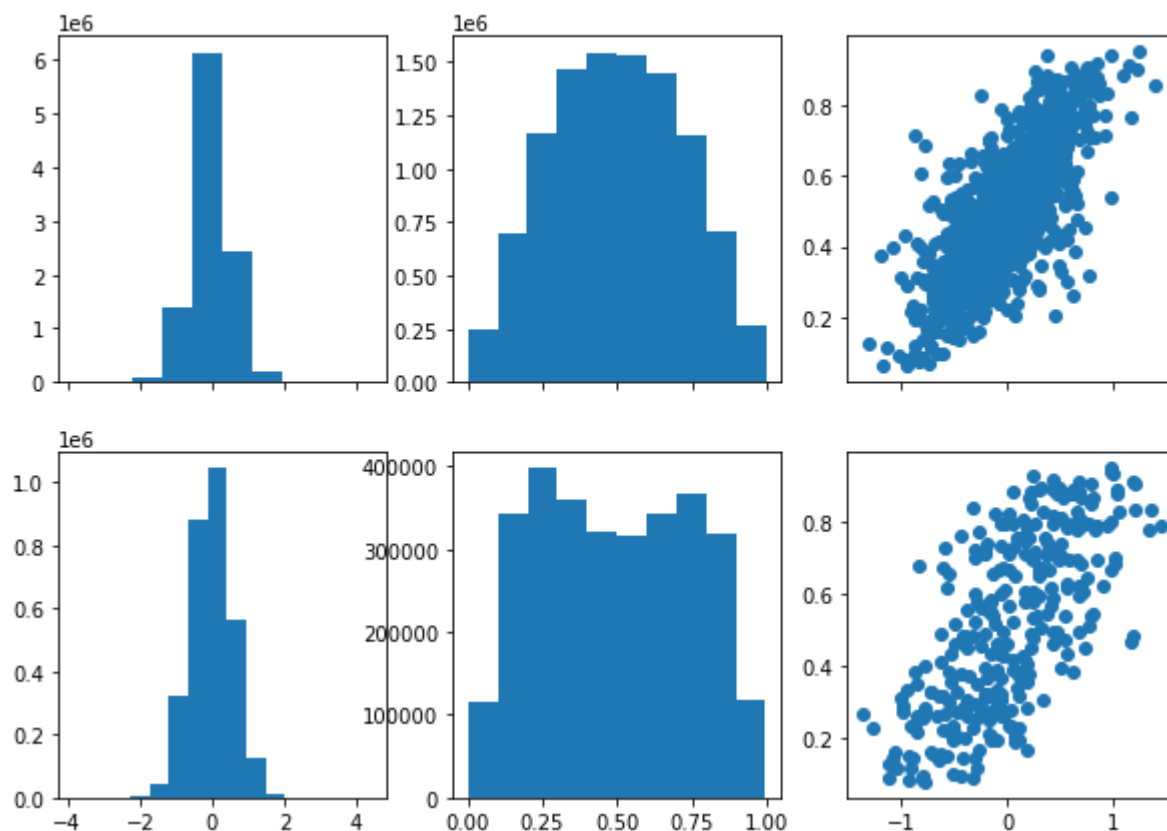
Batch Normalization

22



```
1 class MLP_SigmoidBatchNorm(nn.Module):
2     def __init__(self, in_features, out_features):
3         super(MLP_SigmoidBatchNorm, self).__init__()
4
5         self.hidden_layer1 = 1024
6         self.hidden_layer2 = 300
7
8         self.fc1 = nn.Linear(in_features, self.hidden_layer1)
9         self.bn1 = nn.BatchNorm1d(self.hidden_layer1)
10        self.act1 = nn.Sigmoid()
11        self.fc2 = nn.Linear(self.hidden_layer1, self.hidden_layer2)
12        self.bn2 = nn.BatchNorm1d(self.hidden_layer2)
13        self.act2 = nn.Sigmoid()
14        self.fc3 = nn.Linear(self.hidden_layer2, out_features)
15
16
17    def forward(self, x):
18        output1 = self.fc1(x)
19        bn1 = self.bn1(output1)
20        hidden1 = self.act1(bn1)
21
22        output2 = self.fc2(hidden1)
23        bn2 = self.bn2(output2)
24        hidden2 = self.act2(bn2)
25
26        onehot_out = self.fc3(hidden2)
27
28        return output1, hidden1, output2, hidden2, onehot_out
29
30 init_model(MLP_SigmoidBatchNorm(len(train_0_x.view([-1])), 10).to(device))
31 plot_inner_dist()
```

Epoch: 10 / Accuracy: 0.9645999670028687




Batch Normalization


23



Epoch: 10 / Accuracy: 0.9372999668121338



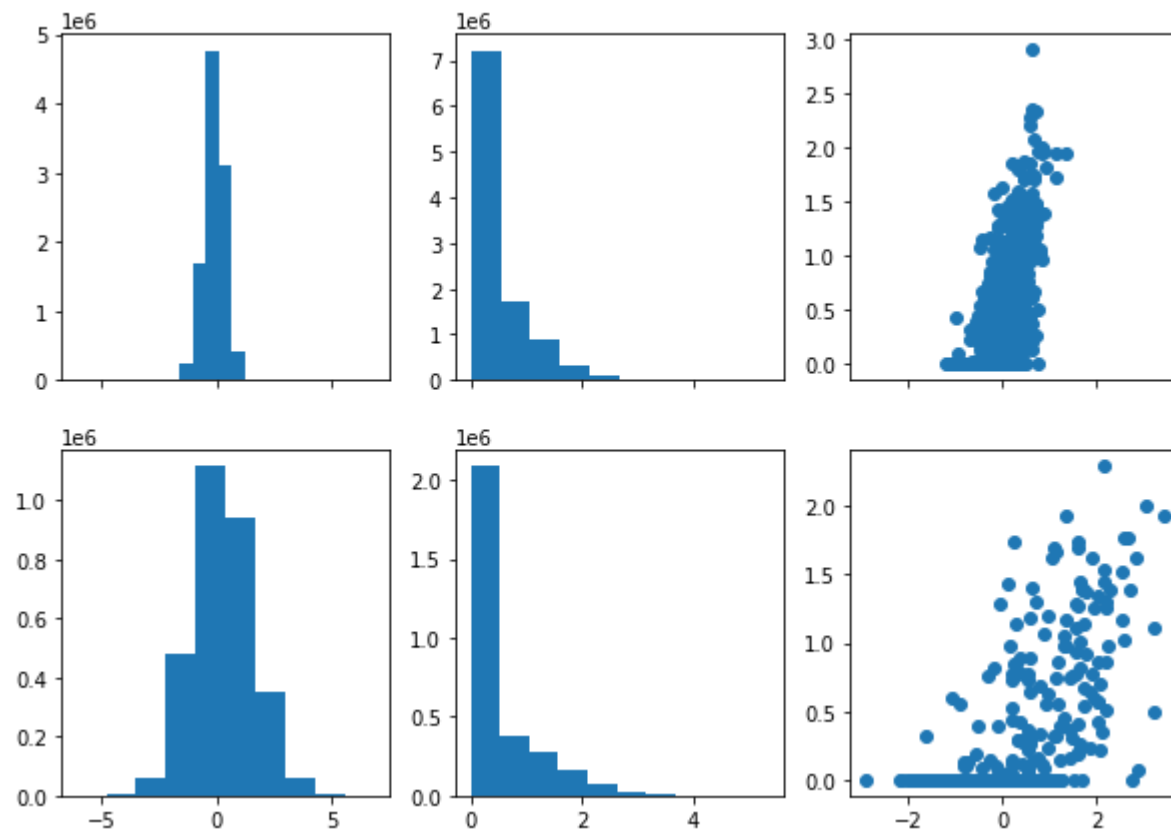
Epoch: 10 / Accuracy: 0.9645999670028687





```
1 class MLPReLUBatchNorm(nn.Module):
2     def __init__(self, in_features, out_features):
3         super(MLPReLUBatchNorm, self).__init__()
4
5         self.hidden_layer1 = 1024
6         self.hidden_layer2 = 300
7
8         self.fc1 = nn.Linear(in_features, self.hidden_layer1)
9         self.bn1 = nn.BatchNorm1d(self.hidden_layer1)
10        self.act1 = nn.ReLU()
11        self.fc2 = nn.Linear(self.hidden_layer1, self.hidden_layer2)
12        self.bn2 = nn.BatchNorm1d(self.hidden_layer2)
13        self.act2 = nn.ReLU()
14        self.fc3 = nn.Linear(self.hidden_layer2, out_features)
15
16    def forward(self, x):
17        output1 = self.fc1(x)
18        bn1 = self.bn1(output1)
19        hidden1 = self.act1(bn1)
20
21        output2 = self.fc2(hidden1)
22        bn2 = self.bn2(output2)
23        hidden2 = self.act2(bn2)
24
25        onehot_out = self.fc3(hidden2)
26
27        return output1, hidden1, output2, hidden2, onehot_out
28
29 init_model(MLPReLUBatchNorm(len(train_0_x.view([-1])), 10).to(device))
30 plot_inner_dist()
```

Epoch: 10 / Accuracy: 0.9833999872207642

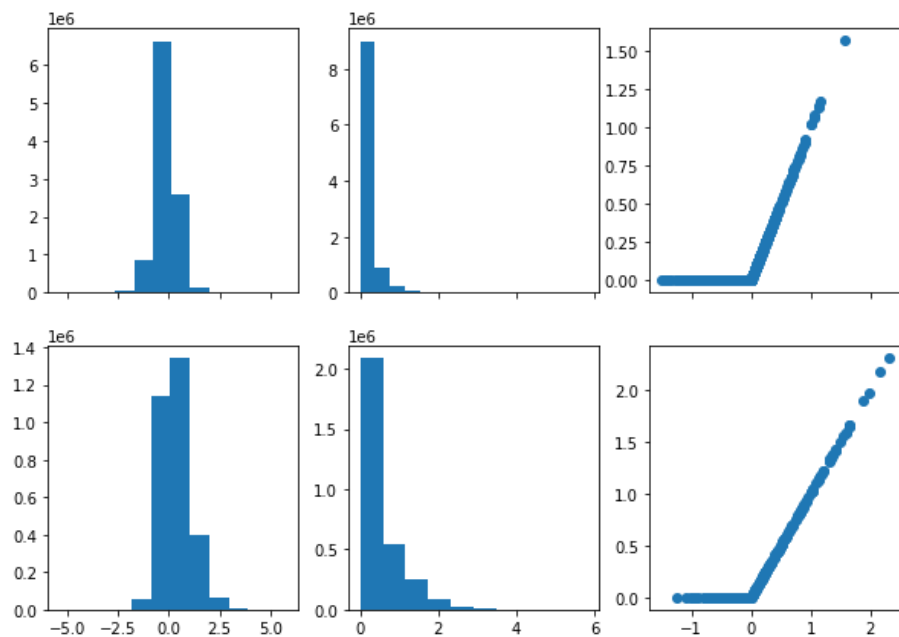


Batch Normalization

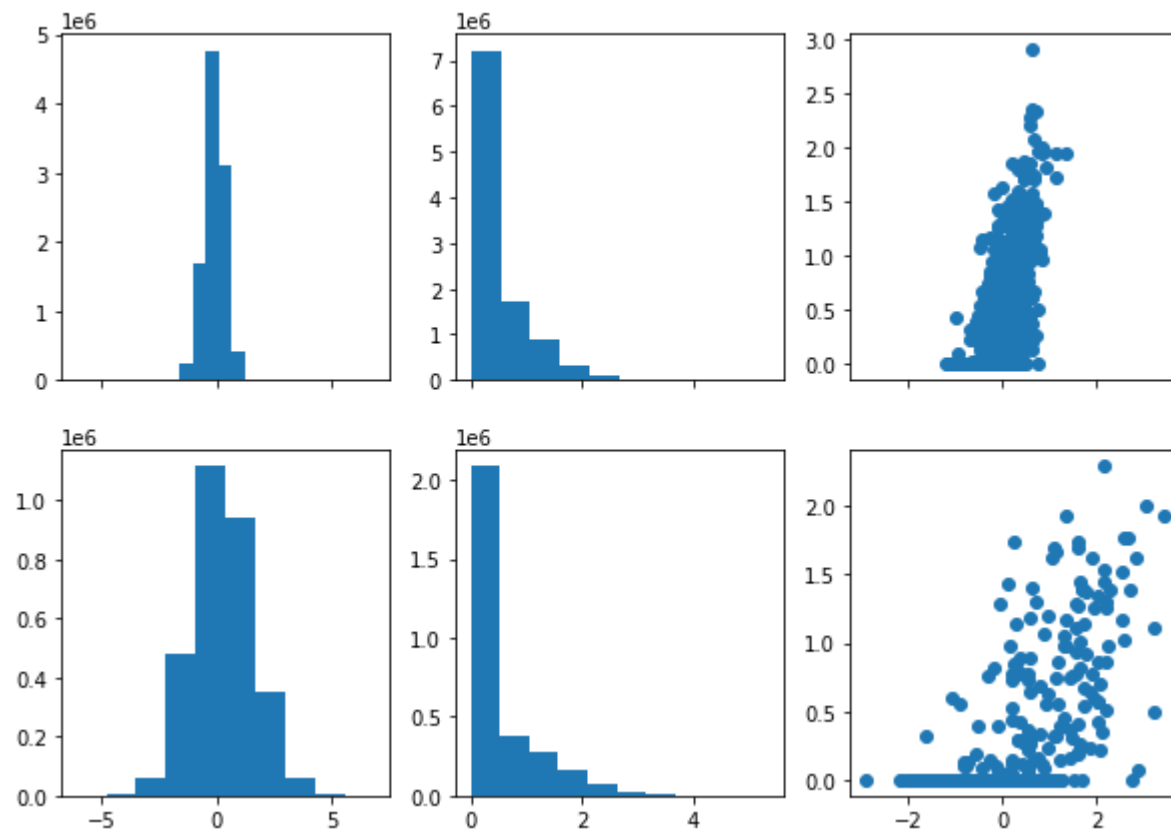
25



Epoch: 10 / Accuracy: 0.9822999835014343





Epoch: 10 / Accuracy: 0.9833999872207642



Dropout



26






○ :units turned off

얼굴위주



○ :units turned off

색지우고



○ :units turned off

귀 빼고






```
1 class Dropout(nn.Module):
2     def __init__(self, in_features, out_features):
3         super(Dropout, self).__init__()
4
5         self.hidden_layer = 32
6         self.dropout_rate = .2 # probability
7
8         self.fc1 = nn.Linear(in_features, self.hidden_layer1)
9         self.act1 = nn.ReLU()
10        self.drop1 = nn.Dropout(dropout_rate)
11        self.fc2 = nn.Linear(self.hidden_layer2, out_features)
12
13    def forward(self, x):
14        hidden1 = self.act1(self.fc1(x))
15        drop1 = self.drop1(hidden1)
16
17        onehot_out = self.fc2(hidden2)
18
19        return onehot_out
```



```
1 # Training Initialization
2 init_model(Dropout(len(train_D_x.view([-1]))), 10))
3 init_epoch()
4 init_log()
5
6 # Training Iteration
7 while epoch_not_finished():
8     start_time = time.time()
9     net.train()
10    tloss, tacc, _, _ = epoch(train_loader)
11    end_time = time.time()
12    time_taken = end_time - start_time
13    record_train_log(tloss, tacc, time_taken)
14    with torch.no_grad():
15        net.eval()
16        vloss, vacc, _, _ = epoch(valid_loader)
17        record_valid_log(vloss, vacc)
18    print_log()
19
20 print('\n Training completed!')
21
22 # Accuracy for test dataset
23 with torch.no_grad():
24     net.eval()
25     test_loss, test_acc, test_out, test_label = epoch(test_loader)
26     print('\n===== Test Result =====\n')
27     print(f'Test accuracy = {test_acc}\nTest loss = {test_loss}')
```



Anything Strange?



===== Test Result =====

Test accuracy = 0.9611999988555908

Test loss = 0.13464893400669098




Make
your own
Model !



```
1 class MLPDropout(nn.Module):
2     def __init__(self, in_features, out_features):
3         super(MLPDropout, self).__init__()
4
5         self.hidden_layer1 = 165
6         self.hidden_layer2 = 165
7         self.dropout_rate = 0.2
8
9         self.layer1 = nn.Sequential(
10             nn.Linear(in_features, self.hidden_layer1),
11             nn.BatchNorm1d(self.hidden_layer1),
12             nn.ReLU(),
13             nn.Dropout(self.dropout_rate)
14         )
15         self.layer2 = nn.Sequential(
16             nn.Linear(self.hidden_layer1, self.hidden_layer2),
17             nn.BatchNorm1d(self.hidden_layer2),
18             nn.ReLU(),
19             nn.Dropout(self.dropout_rate)
20         )
21         self.out_layer = nn.Linear(self.hidden_layer2, out_features)
22
23     def forward(self, x):
24
25         hidden1 = self.layer1(x)
26         hidden2 = self.layer2(hidden1)
27
28         onehot_out = self.out_layer(hidden2)
29
30         return onehot_out
```


Batch Normalization + Dropout

33



Base.ipynb :

<https://git.io/aibd-mlp-3>

Iter: 25 >> T_loss 0.04564 T_acc 0.98463 V_loss 0.04580 V_acc 0.98667 3.829s

Training completed!

===== Test Result =====

Test accuracy = 0.9837999939918518
Test loss = 0.053974274545907974

Better?



Full code :

`https://git.io/aibd-mlp-3-full`