

June 14th, 2021

AI+BD ML Lab. Day 4

Convolutional Neural Network (CNN)

YoungIn Kim
<youngkim21@postech.edu>



Contents

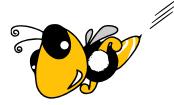
1. Today's Goals
2. Kaggle competition (for test)
3. “CNN”
 - ★ Create CNN model with CAM
 - ★ Learn CIFAR-10 data with your model
 - ★ Watch your CAM
4. Early stopping



1. Understanding **basic components of CNN**
2. Composing a **CNN model** and test it with CIFAR-10 dataset
3. Watching **Class Activation Map** using CNN

Kaggle competition

Let's get the highest score!



- **Link**

- <https://www.kaggle.com/t/c5ba45a8b05443f3b1a62c50cb9188a6>

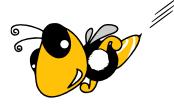
 InClass Prediction Competition

POSCO AI & Big Data Academy (21-2)

Make BEST accuracy for character data

3 days to go

3 days to go



- **Sample code**
 - <https://git.io/kaggle-sample>
 - need to change “root” & “save_root”

Backgrounds

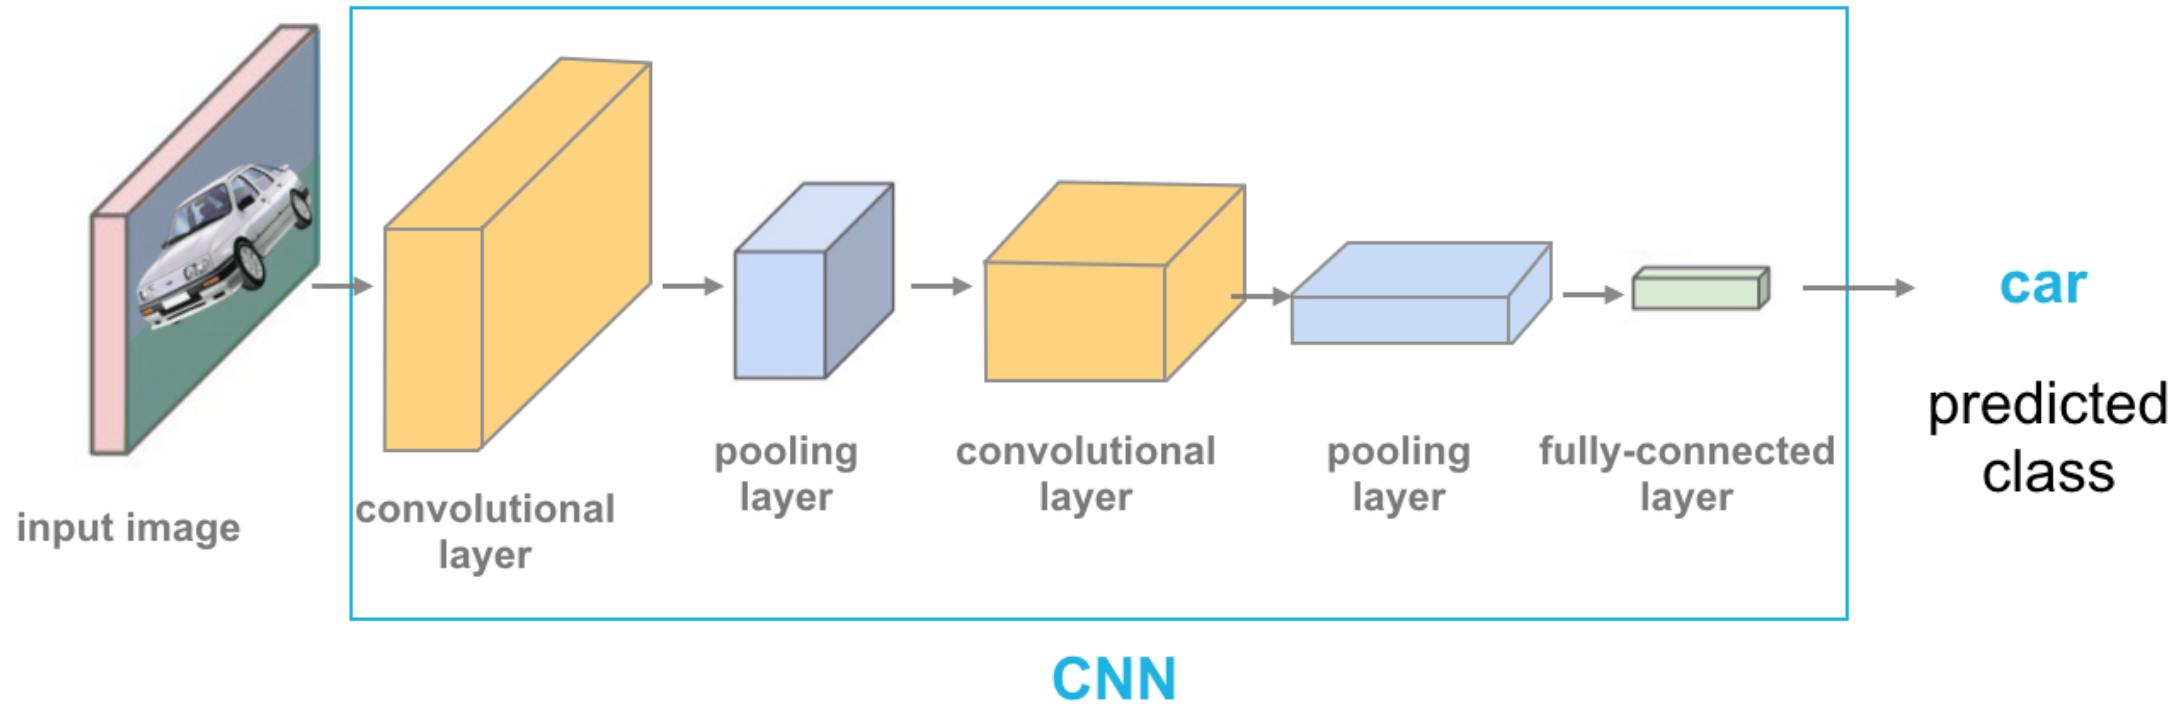
For your information & reminding

Convolutional Neural Network

8



- A good network for image



Components of CNN

9



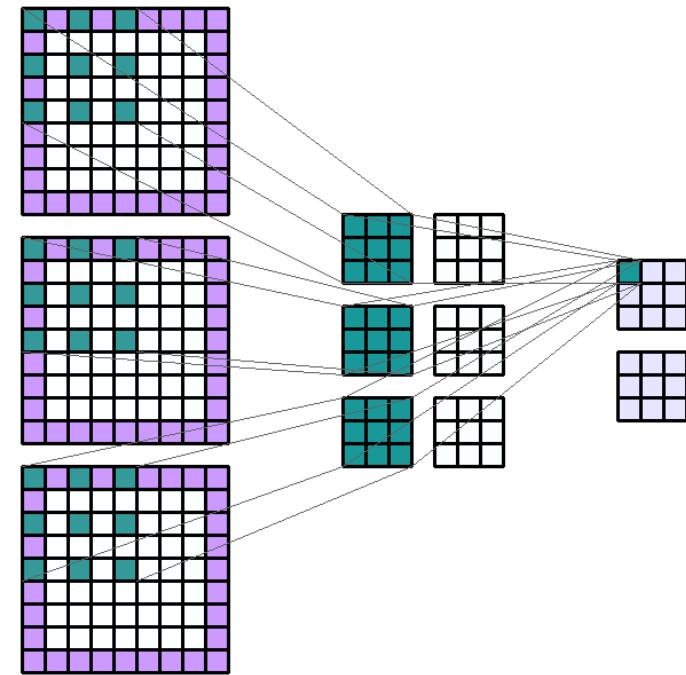
▪ Convolution Operation (Conv2d)

0	0	0	0	0	0	0	0
0	60	113	56	139	85	0	0
0	73	121	54	84	128	0	0
0	131	99	70	129	127	0	0
0	80	57	115	69	134	0	0
0	104	126	123	95	130	0	0
0	0	0	0	0	0	0	0

Kernel

0	-1	0
-1	5	-1
0	-1	0

114			

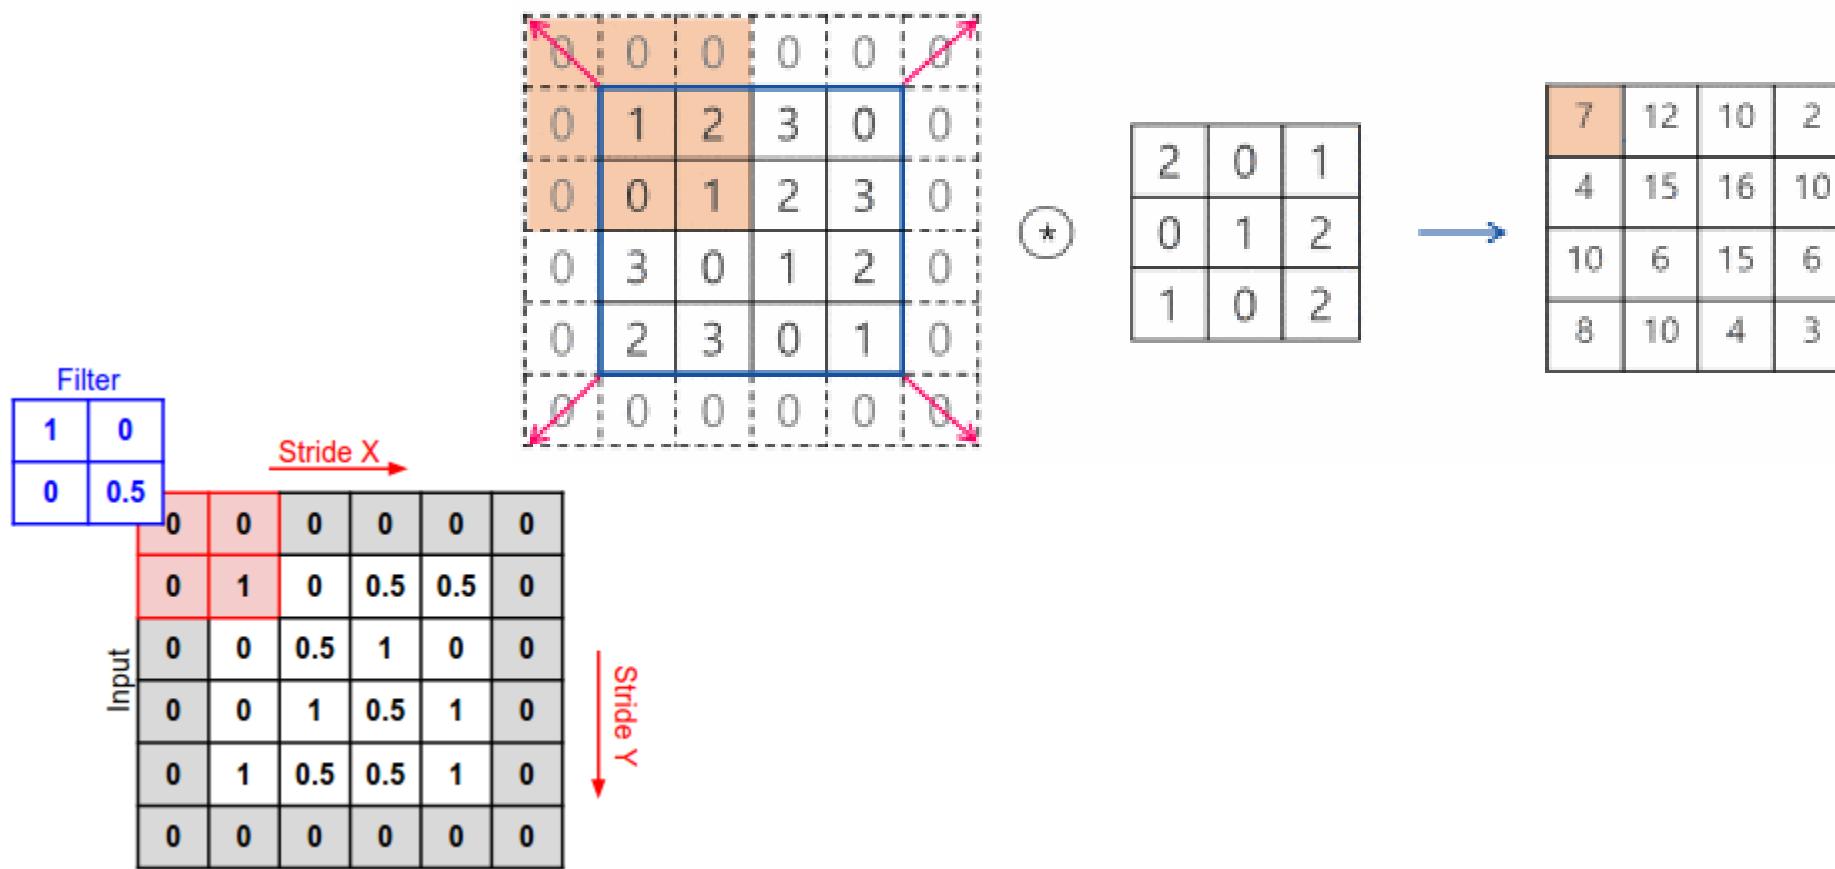


Components of CNN

10

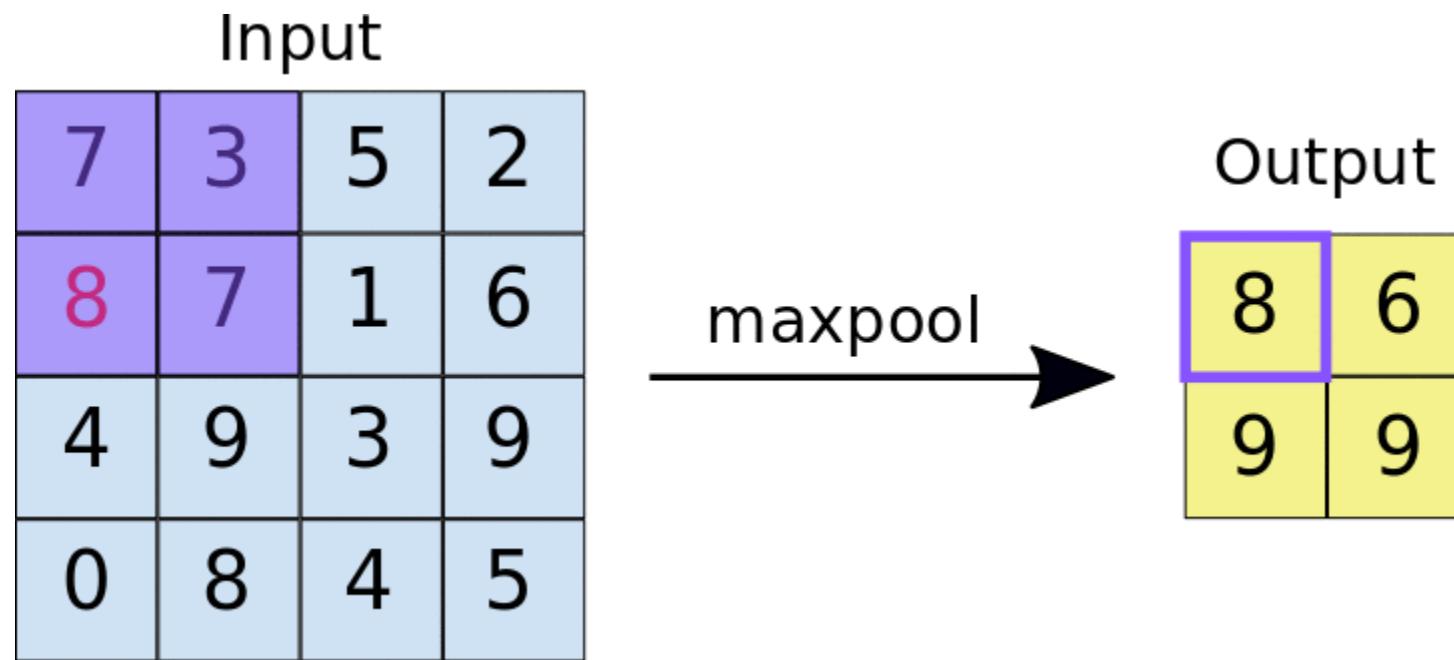


- Convolution Operation (Conv2d) - Stride, Padding





- Max-pooling (MaxPool2d)

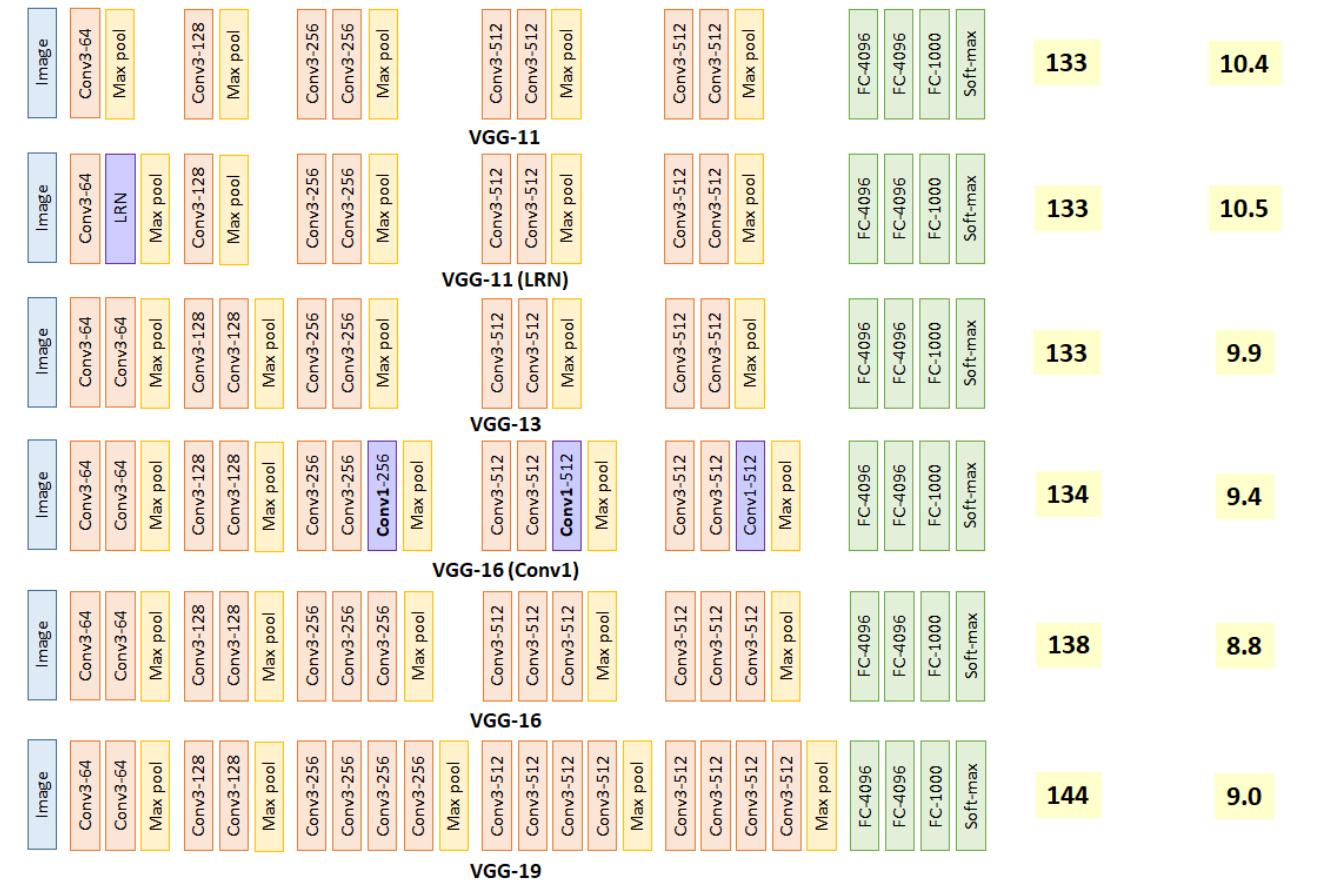
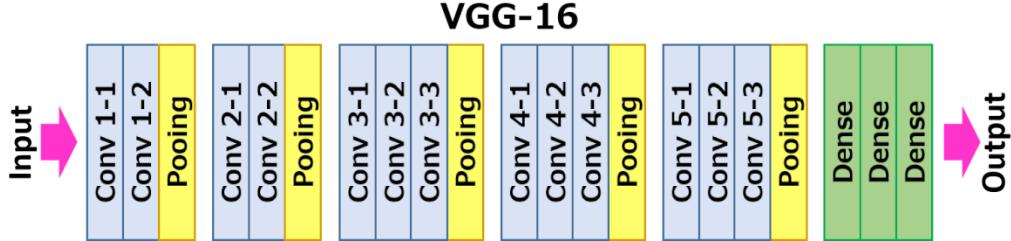
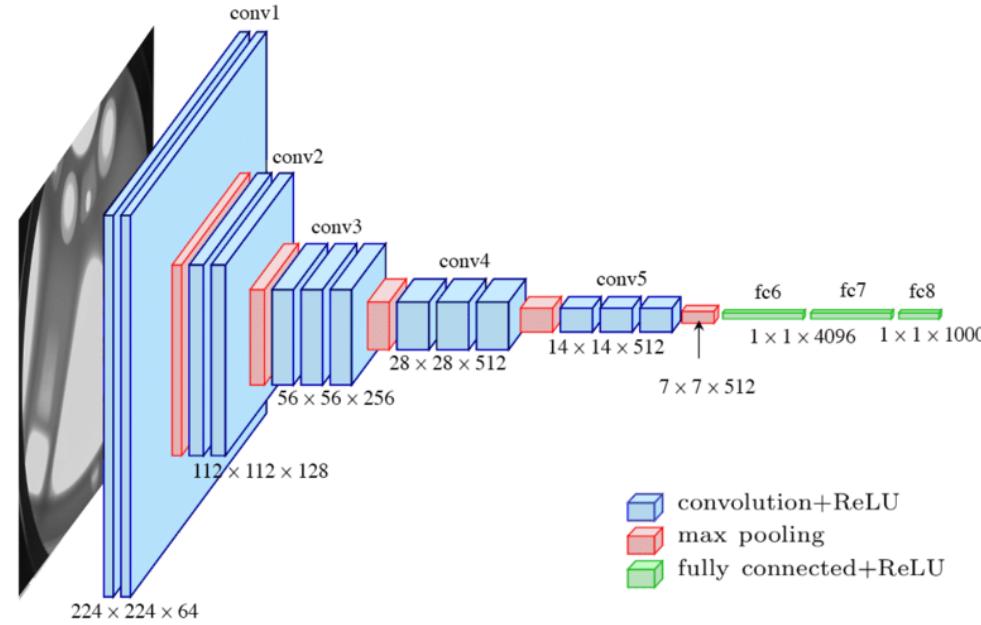


Famous CNNs

12

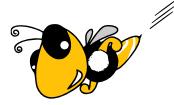


VGG11-19

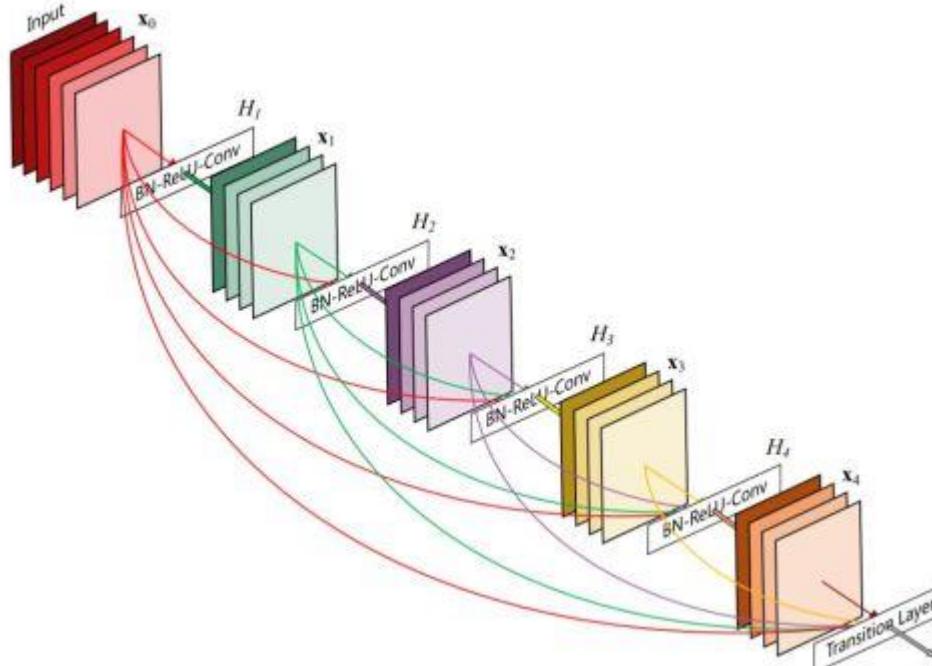
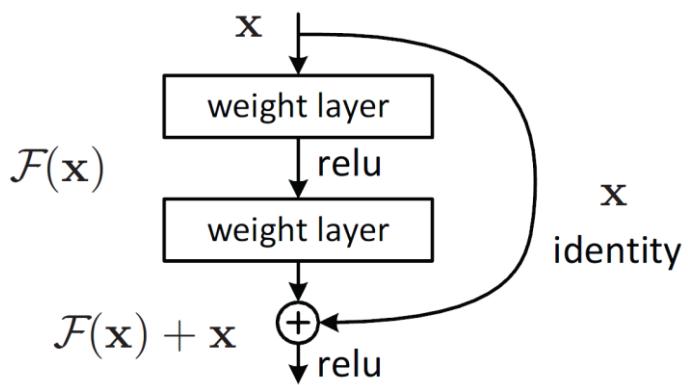


Famous CNNs

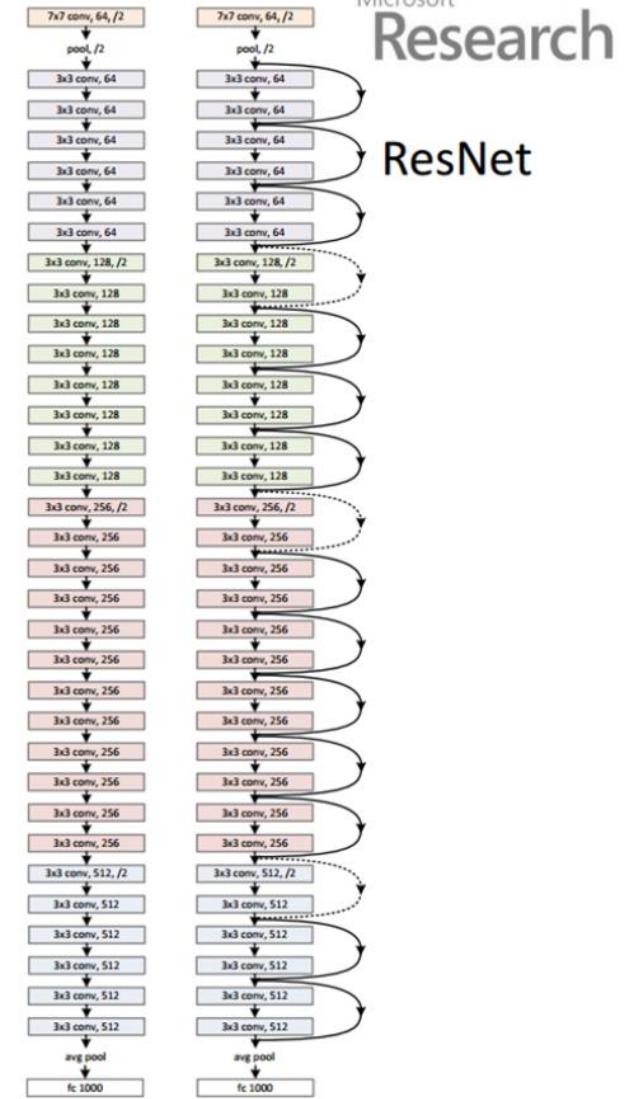
13



ResNet



plain net

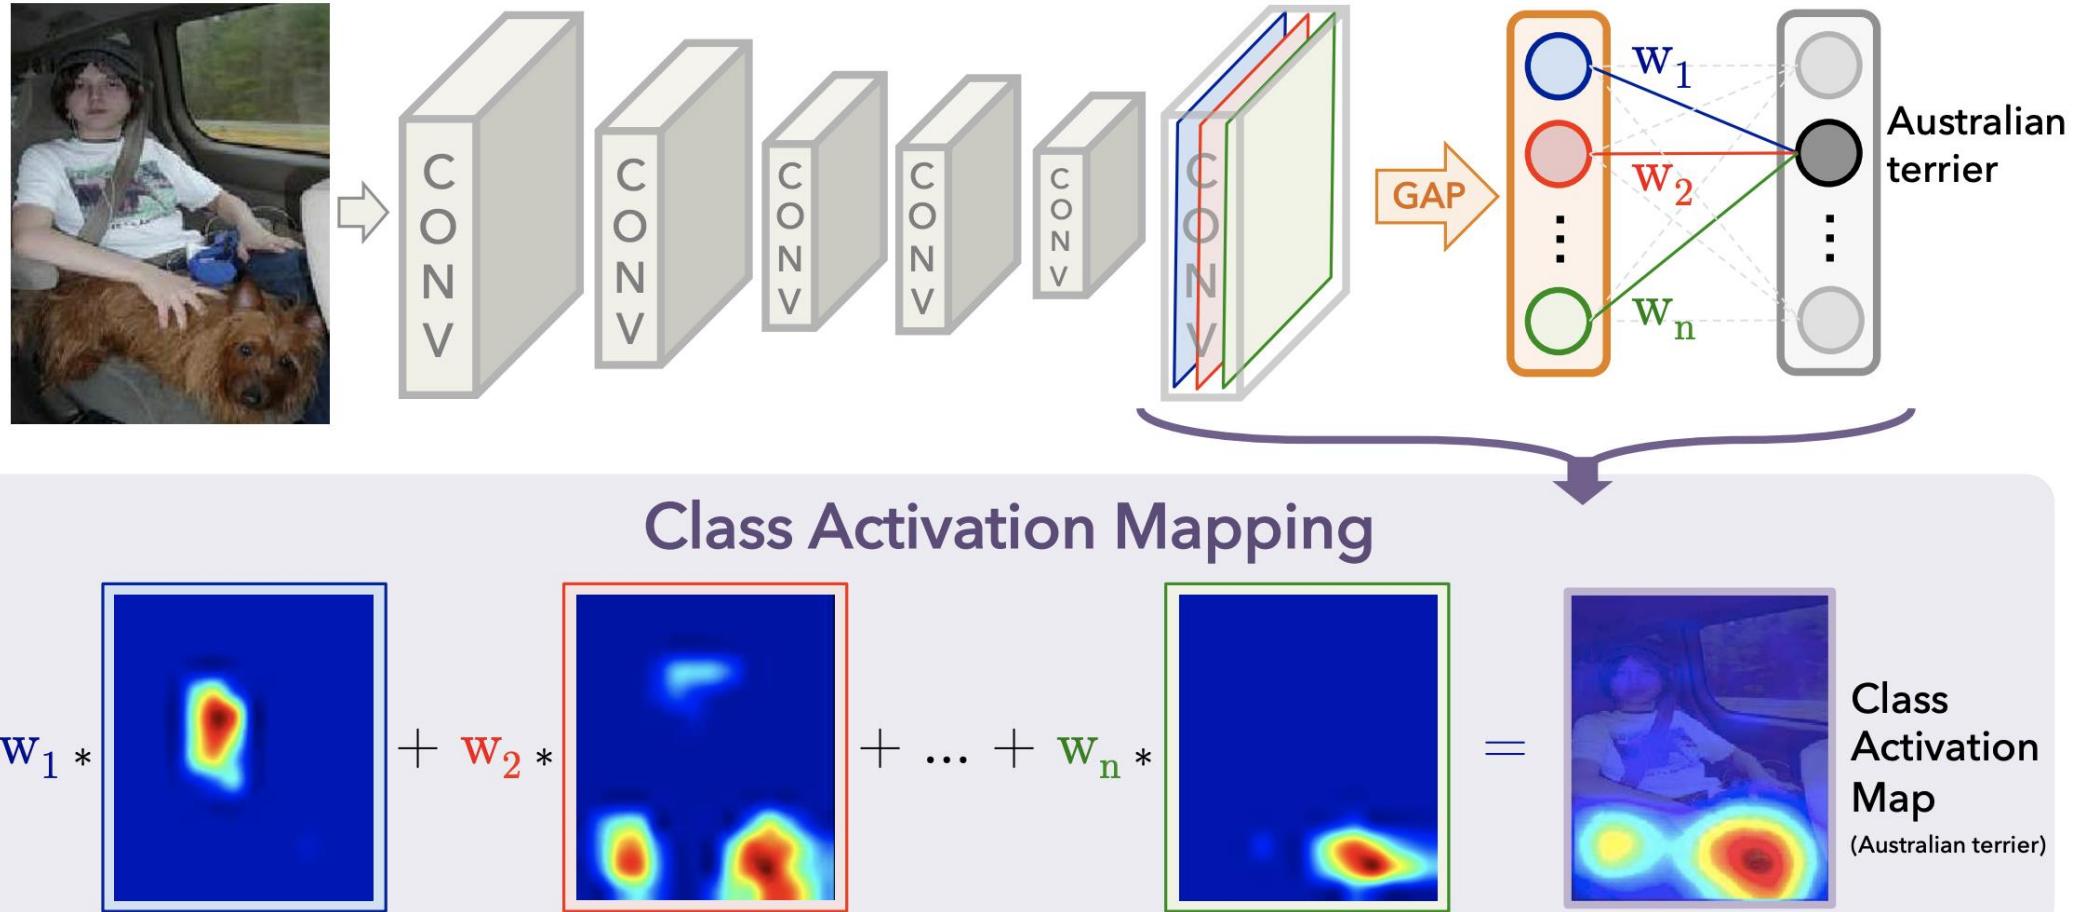


Microsoft
Research

ResNet

Class Activation Map

14





10 classification dataset by Canadian Institute For Advanced Research

비행기



자동차



새



고양이



사슴



개



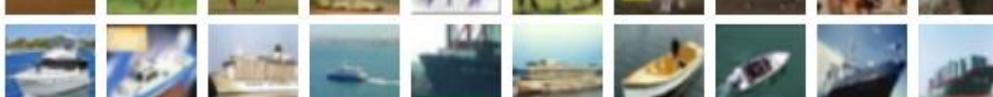
개구리



말



배



트럭



32X32 RGB image

50000 training data pairs

10000 test data pairs

Today we will upscale this
32x32 image to 224x224 RGB
image

It's coding time

Let's fill the I.P.Y.N.B



Base .ipynb :

<https://git.io/aibd-cnn-student>

Required Python Packages :

- matplotlib numpy scikit-learn **scikit-image**
- torch torchvision **torchsummary**

1. Lib & Dev Prep.



Already filled by TA

```
# You don't need to edit this section today.  
import gc  
import matplotlib.pyplot as plt  
import numpy as np  
import time  
import torch  
import torch.nn as nn  
import torchvision.transforms as transforms  
  
from IPython.display import clear_output  
from skimage.transform import resize as OVRResize  
from torch.cuda import memory_allocated, empty_cache  
from torch.optim import Adam  
from torch.utils.data import DataLoader, random_split  
from torchsummary import summary as Summary  
from torchvision.datasets import CIFAR10  
from torchvision.transforms import Compose, ToTensor, Normalize, RandomCrop, RandomHorizontalFlip, \  
    ToPILImage, Resize, Grayscale  
  
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
print(f'{"CPU" if device == "cpu" else "GPU"} will be used in training/validation.')
```

Today, we will use a lot of preprocessing techniques
(torchvision.transforms)

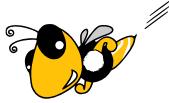


2. Hyperparams

No Model-related Hyper Parameters Today!

```
# Data Load  
input_norm_mean = (.4914, .4822, .4465)  
input_norm_std = (.2023, .1994, .2010)  
batch_size = 16  
  
# Learning  
logging_dispfig = True  
maximum_epoch = 15  
learning_rate = 1e-3
```

- `input_norm_mean` and `input_norm_std` normalizes R/G/B channel to make various images be looked similarly.
- Those values are fit to CIFAR-10, calculated based on computer vision theory.
- Today's batch size is small (16), because of the CUDA memory problem.



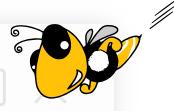
2. Hyperparams

No Model-related Hyper Parameters Today!

```
# Load dataset into python variable
input_transform = transforms.Compose([
    Resize(224),
    RandomCrop(224, padding=24),
    RandomHorizontalFlip(),
    ToTensor(),
    Normalize(mean=input_norm_mean, std=input_norm_std),
])

train_data = CIFAR10("./", train=True, transform=input_transform, download=True)
train_data, valid_data = random_split(train_data, [45000, 5000])
test_data = CIFAR10("./", train=False, transform=input_transform, download=False)
```

- Today, we will use CNN model based on VGG11. VGG11 receives images with size of 224x224, so in same way, we will convert our 32x32 CIFAR10 images into 224x224 images. (“Resize(224)”)
- Then we will apply data augmentation technique for more generic model. (“RandomCrop”, “RandomHorizontalFlip”)
- Then we will convert our data onto Tensor and will apply RGB normalizing as I mentioned in previous slide.



Dataset loading

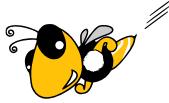
```
# Check the data
print(f'Train dataset length = {len(train_data)}')
print(f'Valid dataset length = {len(valid_data)}')
print(f'Test dataset length = {len(test_data)}\n')

print(f'Classes = {test_data.classes}\n')

train_0_x, train_0_y = train_data[0]
print(f'Label ({type(train_0_y).__name__}) = {train_0_y} ({test_data.classes[train_0_y]})')
print(f'Data ({type(train_0_x).__name__}) = {train_0_x.shape}\n')

fig, (ax1, ax2, ax3) = plt.subplots(1, 3)
fig.suptitle(f'Example data 0 [label={test_data.classes[train_0_y]}]')
ax1.imshow(train_0_x[0, :, :])
ax1.set_title('Ch. 0 (Red)')
ax2.imshow(train_0_x[1, :, :])
ax2.set_title('Ch. 1 (Green)')
ax3.imshow(train_0_x[2, :, :])
ax3.set_title('Ch. 2 (Blue)')
plt.show()
```

- These code will show you R/G/B channel of an example image, in separated.

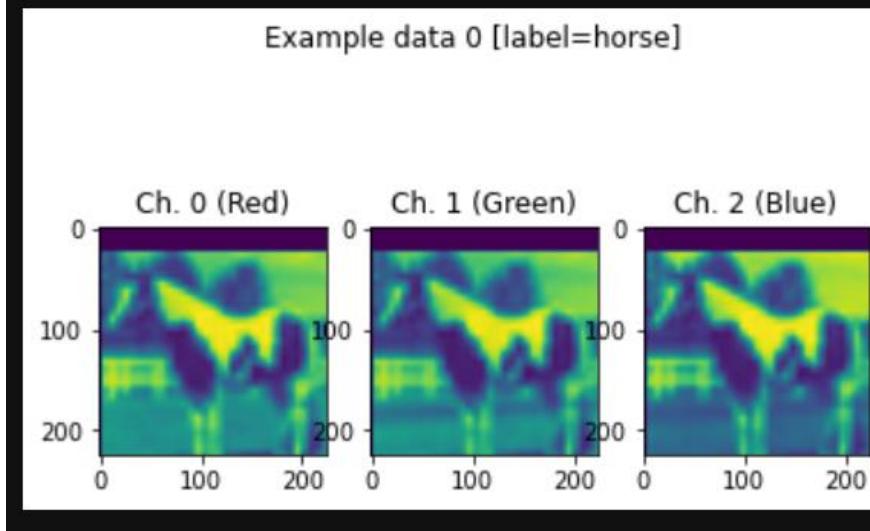


Expected Result

```
Train dataset length = 45000
Valid dataset length = 5000
Test dataset length = 10000

Classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

Label (int) = 7 (horse)
Data (Tensor) = torch.Size([3, 224, 224])
```



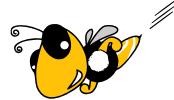
- Look. As we applied RandomCrop, the upper part of the horse image is cropped! (empty section colored as purple).



DataLoader Wrapping

```
# Examine the data loader
train_enumerator = enumerate(train_loader)
ex_batch_idx, (ex_data, ex_label) = next(train_enumerator)
print(f'Idx: {ex_batch_idx} / X.shape = {ex_data.shape} / Y.shape = {ex_label.shape}\n')
print(f'Y[0:{batch_size}] = {ex_label.tolist()}\n')
print(f'→ Label = {np.array(test_data.classes)[ex_label]}\n')

preview_index = 1
fig, (ax1, ax2, ax3) = plt.subplots(1, 3)
fig.suptitle(f'Example data {preview_index} from DataLoader [label={test_data.classes[ex_label[preview_index]]}]')
ax1.imshow(ex_data[preview_index, 0, :, :])
ax1.set_title('Ch. 0 (Red)')
ax2.imshow(ex_data[preview_index, 1, :, :])
ax2.set_title('Ch. 1 (Green)')
ax3.imshow(ex_data[preview_index, 2, :, :])
ax3.set_title('Ch. 2 (Blue)')
plt.show()
```



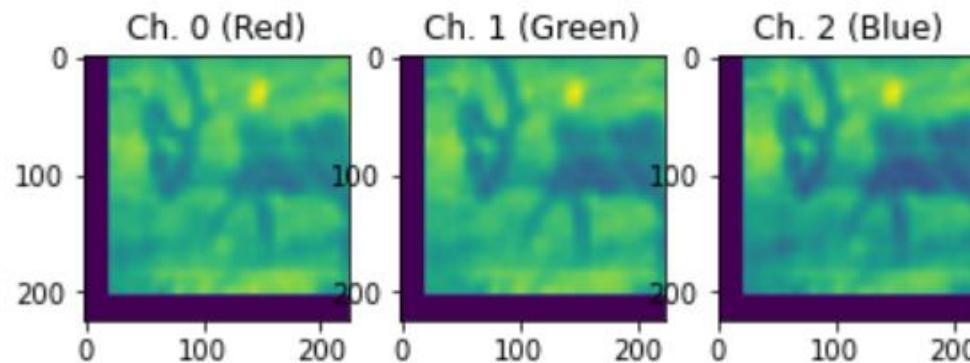
DataLoader Wrapping

Expected Result

```
Idx: 0 / X.shape = torch.Size([16, 3, 224, 224]) / Y.shape = torch.Size([16])

Y[0:16] = [7, 4, 3, 5, 6, 6, 2, 0, 1, 9, 0, 7, 7, 0, 3, 2]
→ Label = ['horse' 'deer' 'cat' 'dog' 'frog' 'frog' 'bird' 'airplane' 'automobile'
'truck' 'airplane' 'horse' 'horse' 'airplane' 'cat' 'bird']
```

Example data 1 from DataLoader [label=deer]

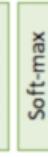
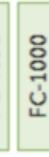
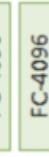
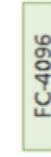
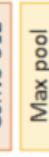
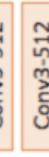
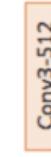
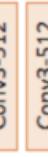
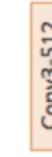
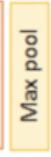
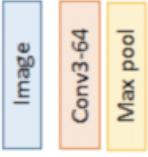




Model Definition

```
- □ ×  
  
def init_model():  
    global net, loss_fn, optim  
    net = MyLittleCNN().to(device)  
    loss_fn = nn.CrossEntropyLoss()  
    optim = Adam(net.parameters(), lr=learning_rate)
```

- As we did yesterday, we will define our model in class structure.
- Loss function is still crossentropy
- Today, we will use “Adam” optimizer, which is widely used these days instead of SGD.



VGG-11

4. Fn. Definitions

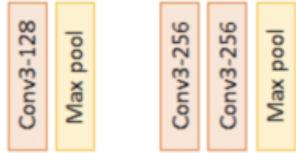
Model Definition

26

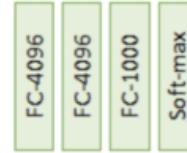


```
def __init__(self):
    super(MyLittleCNN, self).__init__()
    self.convolution_part = nn.Sequential(
        nn.Conv2d(3, 64, kernel_size=3, padding=1), nn.BatchNorm2d(64), nn.LeakyReLU(0.2),
        nn.MaxPool2d(kernel_size=2, stride=2),
        nn.Conv2d(64, 128, kernel_size=3, padding=1), nn.BatchNorm2d(128), nn.LeakyReLU(0.2),
        nn.MaxPool2d(kernel_size=2, stride=2),
        nn.Conv2d(128, 256, kernel_size=3, padding=1), nn.BatchNorm2d(256), nn.LeakyReLU(0.2),
        nn.Conv2d(256, 256, kernel_size=3, padding=1), nn.BatchNorm2d(256), nn.LeakyReLU(0.2),
        nn.MaxPool2d(kernel_size=2, stride=2),
        nn.Conv2d(256, 512, kernel_size=3, padding=1), nn.BatchNorm2d(512), nn.LeakyReLU(0.2),
        nn.Conv2d(512, 512, kernel_size=3, padding=1), nn.BatchNorm2d(512), nn.LeakyReLU(0.2),
        nn.MaxPool2d(kernel_size=2, stride=2),
        nn.Conv2d(512, 512, kernel_size=3, padding=1), nn.BatchNorm2d(512), nn.LeakyReLU(0.2),
        nn.Conv2d(512, 512, kernel_size=3, padding=1), nn.BatchNorm2d(512), nn.LeakyReLU(0.2),
        nn.MaxPool2d(kernel_size=2, stride=2)
    )
    self.channelavg_part = nn.AvgPool2d(7)
    self.classifier_part = nn.Linear(512, 10, bias=False)
```

- Compare top-left VGG11 image and above code.



VGG-11



4. Fn. Definitions

Model Definition

27



```
def forward(self, data):
    conv_out = self.convolution_part(data)
    avg_out = self.channelavg_part(conv_out)
    avg_out_flatten = avg_out.reshape(avg_out.size(0), -1)
    classifier_out = self.classifier_part(avg_out_flatten)
    return classifier_out, conv_out
```

- Forward function only connects “convolution_part -> channelavg_part -> classifier_part”, with reshaping output of channelavg part and taking away convolution part output for purpose of drawing CAM.

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 224, 224]	1,792
BatchNorm2d-2	[-1, 64, 224, 224]	128
LeakyReLU-3	[-1, 64, 224, 224]	0
MaxPool2d-4	[-1, 64, 112, 112]	0
Conv2d-5	[-1, 128, 112, 112]	73,856
BatchNorm2d-6	[-1, 128, 112, 112]	256
LeakyReLU-7	[-1, 128, 112, 112]	0
MaxPool2d-8	[-1, 128, 56, 56]	0
Conv2d-9	[-1, 256, 56, 56]	295,168
BatchNorm2d-10	[-1, 256, 56, 56]	512
LeakyReLU-11	[-1, 256, 56, 56]	0
Conv2d-12	[-1, 256, 56, 56]	590,080
BatchNorm2d-13	[-1, 256, 56, 56]	512
LeakyReLU-14	[-1, 256, 56, 56]	0
MaxPool2d-15	[-1, 256, 28, 28]	0
Conv2d-16	[-1, 512, 28, 28]	1,180,160
BatchNorm2d-17	[-1, 512, 28, 28]	1,024
LeakyReLU-18	[-1, 512, 28, 28]	0
Conv2d-19	[-1, 512, 28, 28]	2,359,808
BatchNorm2d-20	[-1, 512, 28, 28]	1,024
LeakyReLU-21	[-1, 512, 28, 28]	0
MaxPool2d-22	[-1, 512, 14, 14]	0
Conv2d-23	[-1, 512, 14, 14]	2,359,808
BatchNorm2d-24	[-1, 512, 14, 14]	1,024
LeakyReLU-25	[-1, 512, 14, 14]	0
Conv2d-26	[-1, 512, 14, 14]	2,359,808
BatchNorm2d-27	[-1, 512, 14, 14]	1,024
LeakyReLU-28	[-1, 512, 14, 14]	0
MaxPool2d-29	[-1, 512, 7, 7]	0
AvgPool2d-30	[-1, 512, 1, 1]	0
Linear-31	[-1, 10]	5,120

Total params: 9,231,104

Trainable params: 9,231,104

Non-trainable params: 0

Input size (MB): 0.57

Forward/backward pass size (MB): 181.65

Params size (MB): 35.21

Estimated Total Size (MB): 217.44

4. Fn. Definitions

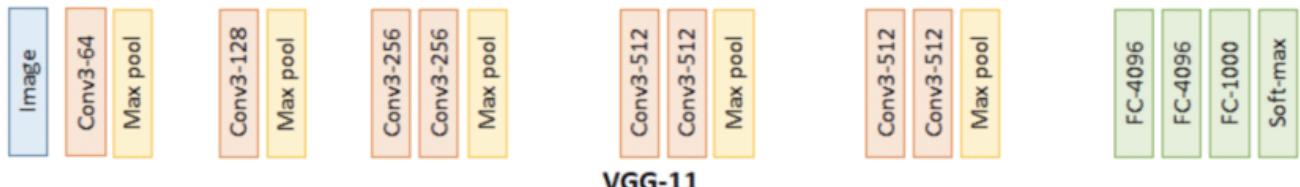
Model Definition

28



Expected Result

- Compare left TorchSummary output and below VGG 11 block diagram.





GPU Memory Cleaner

— □ ×

```
# Memory cleaner to prevent CUDA out of memory error
def clear_memory():
    if device != 'cpu':
        empty_cache()
    gc.collect()
```

- Today's model is based on VGG and VGG is very memory-intensive model. So we're defining GPU/CPU useless cache memory cleaning function.

```

def init_epoch():
    global epoch_cnt
    epoch_cnt = 0

def epoch(data_loader):
    # One epoch : gets data_loader as input and returns loss / accuracy, and
    #             last prediction value / its label(truth) value for future use
    global epoch_cnt
    iter_loss, iter_acc = [], []
    last_grad_performed = False

    # Mini-batch iterations
    for _data, _label in data_loader:
        data, label = _data.to(device), _label.to(device)

        # 1. Feed-forward
        onehot_out, _ = net(data)

        # 2. Calculate accuracy
        _, out = torch.max(onehot_out, 1)
        acc_partial = (out == label).float().sum()
        acc_partial = acc_partial / len(label)
        iter_acc.append(acc_partial.item())

        # 3. Calculate loss
        loss = loss_fn(onehot_out, label)
        iter_loss.append(loss.item())

        # 4. Backward propagation if not in `torch.no_grad()`
        if onehot_out.requires_grad:
            optim.zero_grad()
            loss.backward()
            optim.step()
            last_grad_performed = True

    # Up epoch count if backward propagation is done
    if last_grad_performed:
        epoch_cnt += 1

    # Clear memory to prevent CUDA memory error
    clear_memory()

    return np.average(iter_loss), np.average(iter_acc)

def epoch_not_finished():
    # For now, let's repeat training fixed times, e.g. 25 times.
    # We will learn how to determine training stop or continue later.
    return epoch_cnt < maximum_epoch

```

4. Fn. Definitions

30



Epoch Function

Filled by TA

- Same as yesterday.



4. Fn. Definitions

Logging Function

Filled by TA

```
# Logging (same with day 2)
def init_log():
    global log_stack, iter_log, tloss_log, tacc_log, vloss_log, vacc_log, time_log
    iter_log, tloss_log, tacc_log, vloss_log, vacc_log = [], [], [], [], []
    time_log, log_stack = [], []

def record_train_log(_tloss, _tacc, _time):
    # Push time, training loss, training accuracy, and epoch count into lists
    time_log.append(_time)
    tloss_log.append(_tloss)
    tacc_log.append(_tacc)
    iter_log.append(epoch_cnt)

def record_valid_log(_vloss, _vacc):
    # Push validation loss and validation accuracy into each list
    vloss_log.append(_vloss)
    vacc_log.append(_vacc)

def last(log_list):
    # Get the last member of list. If empty, return -1.
    if len(log_list) > 0: return log_list[len(log_list) - 1]
    else: return -1

def print_log():
    # Generate log string and put it into log stack
    log_str = f'Iter: {last(iter_log):>4d} >> T_loss {last(tloss_log):<8.5f}  '\
              + f'T_acc {last(tacc_log):<6.5f}  V_loss {last(vloss_log):<8.5f}  '\
              + f'V_acc {last(vacc_log):<6.5f}  Ⓢ {last(time_log):5.3f}s'
    log_stack.append(log_str)

# Draw figure if want
if logging_dispfig:
    hist_fig, loss_axis = plt.subplots(figsize=(10, 3), dpi=99)
    hist_fig.patch.set_facecolor('white')

# Draw loss lines
loss_t_line = plt.plot(iter_log, tloss_log, label='Train Loss', color='#FF9999', marker='o')
loss_v_line = plt.plot(iter_log, vloss_log, label='Valid Loss', color='#99B0FF', marker='s')
loss_axis.set_xlabel('epoch')
loss_axis.set_ylabel('loss')

# Draw accuracy lines
acc_axis = loss_axis.twinx()
acc_t_line = acc_axis.plot(iter_log, tacc_log, label='Train Acc.', color='#FF0000', marker='+')
acc_v_line = acc_axis.plot(iter_log, vacc_log, label='Valid Acc.', color='#003AFF', marker='x')
acc_axis.set_ylabel('accuracy')

# Append annotations
hist_lines = loss_t_line + loss_v_line + acc_t_line + acc_v_line
loss_axis.legend(hist_lines, [l.get_label() for l in hist_lines])
loss_axis.grid()
plt.title(f'Learning history until epoch {last(iter_log)}')
plt.draw()

# Print log
clear_output(wait=True)
if logging_dispfig: plt.show()
for idx in reversed(range(len(log_stack))):
    print(log_stack[idx])
```

- Same as yesterday.



5. Training Iter.

Yeah

```
# Training Initialization (same with day 2)
init_model()
init_epoch()
init_log()

# Training Iteration (similar to day 2)
while epoch_not_finished():
    start_time = time.time()
    tloss, tacc = epoch(train_loader)
    end_time = time.time()
    time_taken = end_time - start_time
    record_train_log(tloss, tacc, time_taken)
    with torch.no_grad():
        vloss, vacc = epoch(valid_loader)
        record_valid_log(vloss, vacc)
    print_log()

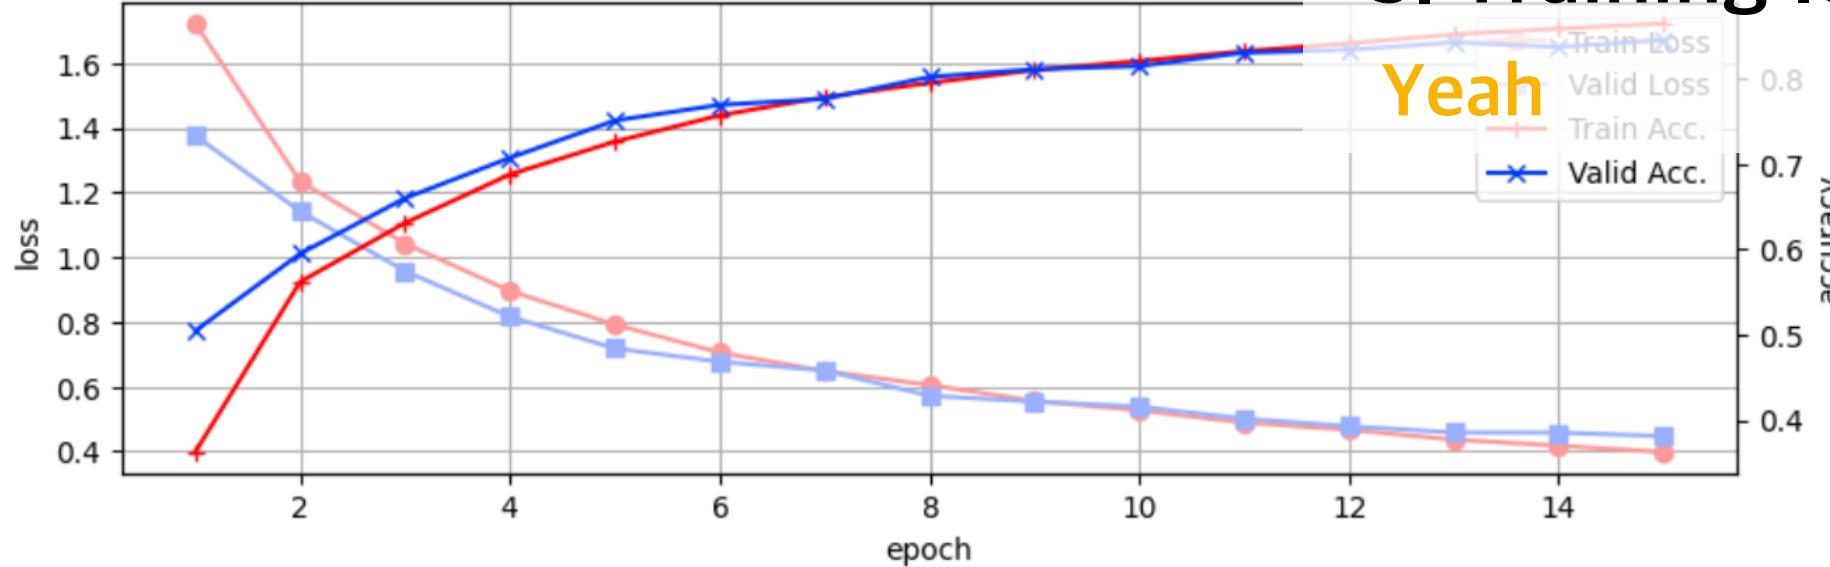
print('\n Training completed!')
```

Filled by TA

- Same as yesterday.



5. Training Iter.



```
Iter: 15 >> T_loss 0.39752    T_acc 0.86513    V_loss 0.44669    V_acc 0.84565    96.090s
Iter: 14 >> T_loss 0.41736    T_acc 0.85895    V_loss 0.45768    V_acc 0.83746    96.118s
Iter: 13 >> T_loss 0.43601    T_acc 0.85220    V_loss 0.45802    V_acc 0.84285    96.044s
Iter: 12 >> T_loss 0.46626    T_acc 0.84173    V_loss 0.47743    V_acc 0.83407    96.128s
Iter: 11 >> T_loss 0.48864    T_acc 0.83253    V_loss 0.50081    V_acc 0.83087    96.179s
Iter: 10 >> T_loss 0.52613    T_acc 0.82101    V_loss 0.53792    V_acc 0.81510    95.971s
Iter: 9 >> T_loss 0.55588     T_acc 0.81041    V_loss 0.55347    V_acc 0.81130    96.406s
Iter: 8 >> T_loss 0.60472     T_acc 0.79523    V_loss 0.57173    V_acc 0.80232    96.224s
Iter: 7 >> T_loss 0.64848     T_acc 0.77869    V_loss 0.64989    V_acc 0.77656    96.172s
Iter: 6 >> T_loss 0.70538     T_acc 0.75727    V_loss 0.67705    V_acc 0.76957    96.181s
Iter: 5 >> T_loss 0.79176     T_acc 0.72648    V_loss 0.71801    V_acc 0.75120    96.270s
Iter: 4 >> T_loss 0.89577     T_acc 0.68794    V_loss 0.81597    V_acc 0.70747    96.073s
Iter: 3 >> T_loss 1.04208     T_acc 0.63111    V_loss 0.95706    V_acc 0.66014    96.143s
Iter: 2 >> T_loss 1.23487     T_acc 0.56203    V_loss 1.14193    V_acc 0.59505    96.072s
Iter: 1 >> T_loss 1.72277     T_acc 0.36240    V_loss 1.37549    V_acc 0.50459    95.061s
```

Training completed!

- Beautiful graph.

Expected Result



Accuracy

- □ ×

```
# Accuracy for test dataset (similar to day 2)
with torch.no_grad():
    test_loss, test_acc = epoch(test_loader)
    print(f'Test accuracy = {test_acc}\nTest loss = {test_loss}')
```

Test accuracy = 0.8406
Test loss = 0.4767277945458889

Expected Result



Class Activation Map

- □ ×

```
# Data preparation (load test data again, but without data augmentation)
cam_data_iterable = CIFAR10(root='./', train=False, download=False)
cam_data = list(cam_data_iterable)
```

- Now we're trying to draw class activation map (CAM).
- Because we want draw the original picture and overlay CAM on it, first load image data without any transformations (first line).
- Then make it as a python list (second line)



Class Activation Map

For the full picture of “Drawing CAM” code,
Look at here : <https://i.imgur.com/tlhLuS7.png>
Also check for “#comments” the description.



Class Activation Map

- □ ×

```
# Draw CAMs
cam_figidx = 1
cam_random_order = np.arange(len(cam_data), dtype=int)
np.random.shuffle(cam_random_order)

fig = plt.figure(figsize=(23, 47))
for class_idx, class_label in enumerate(cam_data_iterable.classes):
    # We will draw 5 CAMs for each classes
    class_axes = [
        fig.add_subplot(10, 5, cam_figidx),
        fig.add_subplot(10, 5, cam_figidx + 1),
        fig.add_subplot(10, 5, cam_figidx + 2),
        fig.add_subplot(10, 5, cam_figidx + 3),
        fig.add_subplot(10, 5, cam_figidx + 4)
    ]
    cam_figidx = cam_figidx + 5

# Find 5 proper data with a specified "label"
class_figidx = 0
```

6. Result Analysis

38



Class Activation Map

```
# We will draw 5 CAMs for each classes
class_axes = [
    fig.add_subplot(10, 5, cam_figidx),
    fig.add_subplot(10, 5, cam_figidx + 1),
    fig.add_subplot(10, 5, cam_figidx + 2),
    fig.add_subplot(10, 5, cam_figidx + 3),
    fig.add_subplot(10, 5, cam_figidx + 4)
]
cam_figidx = cam_figidx + 5

# Find 5 proper data with a specified "label"
class_figidx = 0
for random_order in cam_random_order:
    data, label = cam_data[random_order]
    if label == class_idx:
        # 1. Preprocess original image to put it into the model
        temp_input = ToTensor()(Resize(224)(data))
        temp_input = Normalize(mean=input_norm_mean, std=input_norm_std)(temp_input)
        temp_input = temp_input.reshape((1, 3, 224, 224)).to(device)

        # 2. Put the data into model and get convolutional part output and classifier output
        with torch.no_grad():
            temp_classifier_out, temp_conv_out = net(temp_input)
            _, temp_predicted_class = torch.max(temp_classifier_out, 1)

        # 3. If the predicted label is same with the truth, draw CAM of it
        if temp_predicted_class == label:
            copy_class_weight = list(net.parameters())[1].squeeze()[label]
```

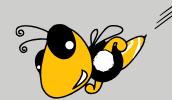
```
temp_input = temp_input.reshape((1, 3, 224, 224)).to(device)
```

```
# 2. Put the data into model and get convolutional part for output  
with torch.no_grad():  
    temp_classifier_out, temp_conv_out = net(temp_input)  
    _, temp_predicted_class = torch.max(temp_classifier_out, 1)
```

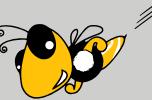
```
# 3. If the predicted label is same with the truth, draw CAM of it
```

```
if temp_predicted_class == label:  
    conv_class_weight = list(net.parameters())[-1].squeeze()[label]  
    temp_conv_out = temp_conv_out.reshape(512, 49)  
    cam = conv_class_weight.matmul(temp_conv_out).reshape(7, 7).cpu().data.numpy()  
    cam = cam - np.min(cam)  
    cam = cam / np.max(cam)  
    class_axes[class_figidx].imshow(  
        Grayscale()(Resize((512, 512), interpolation=0)(data)),  
        cmap='gray',  
    )  
    class_axes[class_figidx].imshow(OVResize(cam, [512, 512]), cmap='jet', alpha=0.4)  
    class_axes[class_figidx].set_xticks([], minor=False)  
    class_axes[class_figidx].set_yticks([], minor=False)  
    if class_figidx == 2:  
        class_axes[class_figidx].set_title(f'\n===== [ {class_label} ] =====\n')  
    class_figidx = class_figidx + 1  
    if class_figidx >= 5:  
        break
```

6. Result Analysis 39



Class Activation Map



6. Result Analysis

40

Class Activation Map

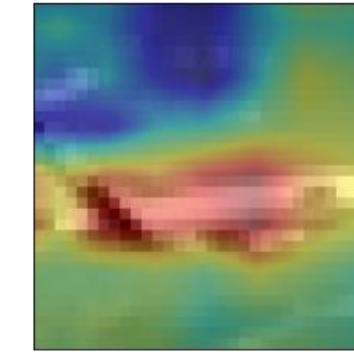
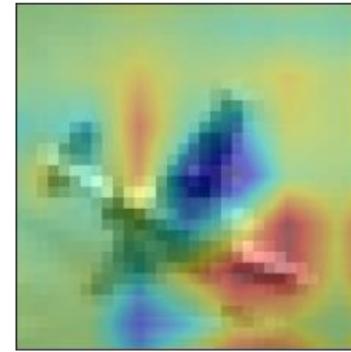
```
temp_predicted_class == label:
    conv_class_weight = list(net.parameters())[-1].squeeze()[label]
    temp_conv_out = temp_conv_out.reshape(512, 49)
    cam = conv_class_weight.matmul(temp_conv_out).reshape(7, 7).cpu().data.numpy()
    cam = cam - np.min(cam)
    cam = cam / np.max(cam)
    class_axes[class_figidx].imshow(
        Grayscale()(Resize((512, 512), interpolation=0)(data)),
        cmap='gray',
    )
    class_axes[class_figidx].imshow(OVResize(cam, [512, 512]), cmap='jet', alpha=0.4)
    class_axes[class_figidx].set_xticks([], minor=False)
    class_axes[class_figidx].set_yticks([], minor=False)
    if class_figidx == 2:
        class_axes[class_figidx].set_title(f'\n===== [ {class_label} ] =====\n')
    class_figidx = class_figidx + 1
    if class_figidx >= 5:
        break

# Tiny error message
if class_figidx != 5:
    print(f'Failed to find proper data for {class_label}[{class_idx}].')
```

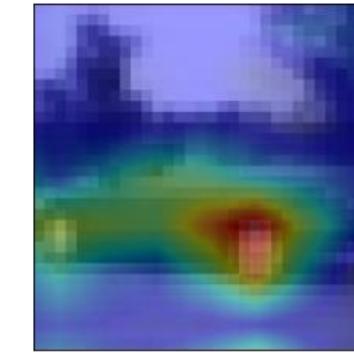
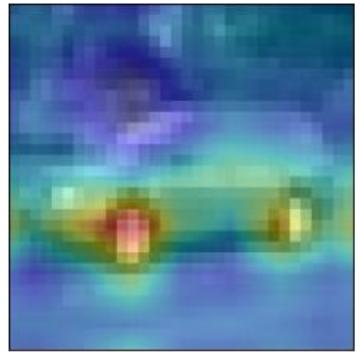
plt.show()



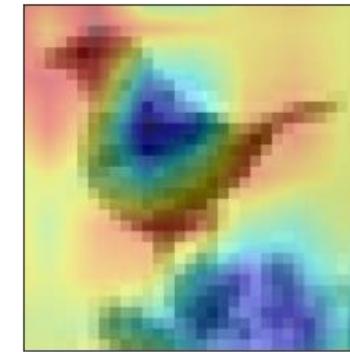
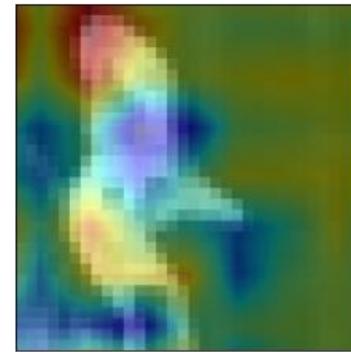
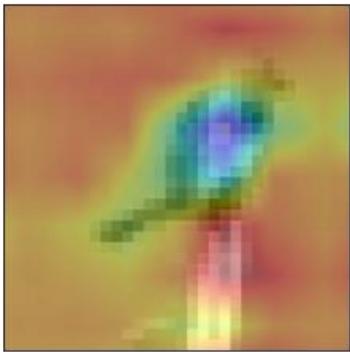
===== [airplane] =====



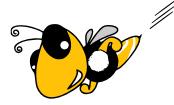
===== [automobile] =====



===== [bird] =====



Expected
Result

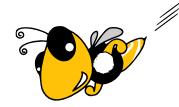


You can check the full version of today's .ipynb at

<https://git.io/aibd-cnn-full>

Early Stopping

Just see it and code



- When the model is too fit to the training dataset so when the performance for non-training data degrades, we call it “overfitted”.
- So how can we prevent it? One way is “early stopping”.



- Code
 - <https://git.io/aibd-early-stop>

```
# Load dataset into python variable
input_transform = transforms.Compose([
    RandomCrop(32, padding=4),
    RandomHorizontalFlip(),
    ToTensor(),
    Normalize(mean=input_norm_mean, std=input_norm_std),
])

train_data = CIFAR10("./", train=True, transform=input_transform, download=True)
train_data, valid_data = random_split(train_data, [2000, 48000])
valid_data, trash_data = random_split(valid_data, [2000, 46000])
test_data = CIFAR10("./", train=False, transform=input_transform, download=False)
```

Just for testing purpose

- **Earlystopping can be done with two mechanism.**
 - Wait for “patience” times of no validation loss decreasing.
 - If the era of no validation loss decreasing continues for “patience” times, then stop the training.
 - After stopping the training, set the parameters of network as its best parameter (who shown the lowest validation loss).



- Earlystopping can be done with two mechanism.

```
.  
def init_epoch():  
    global epoch_cnt, earlystop_cnt  
    epoch_cnt = 0  
    earlystop_cnt = 0  
  
def epoch_not_finished():  
    return earlystop_cnt < earlystop_patience
```

```
# Earlystopping  
def copy_weights(src, dst):  
    dst.load_state_dict(src.state_dict())
```

```
# Training Iteration  
while epoch_not_finished():  
    start_time = time.time()  
    tloss, tacc = epoch(train_loader)  
    end_time = time.time()  
    time_taken = end_time - start_time  
    record_train_log(tloss, tacc, time_taken)  
    with torch.no_grad():  
        vloss, vacc = epoch(valid_loader)  
        if len(vloss_log) and np.min(vloss_log) < vloss:  
            earlystop_cnt = earlystop_cnt + 1  
        else:  
            earlystop_cnt = 0  
            copy_weights(net, best_net)  
            record_valid_log(vloss, vacc)  
            print_log()  
  
copy_weights(best_net, net)
```



- You can get two advantages from earlystopping
 - You don't have to wait for a long time for meaningless training.
 - You can prevent overfitting.
- But sometimes earlystopped model has a poor performance than more-trained model.
 - So, just use earlystopping when you taking aim about the appropriate epoch count.